GROUP WORK PROJECT # 2
GROUP NUMBER: 12132

MScFE 642: Deep Learning for Finance

FULL LEGAL NAME	LOCATION (COUNTRY)	EMAIL ADDRESS	MARK X FOR ANY NON-CONTRIBUTING MEMBER
Kolapo Adedipe	Nigeria	adcolapsmm9@yahoo.com	
Han Li	China	lh15099558125@gmail.com	
Cynthia Njau	Kenya	cynthia.njau98@gmail.com	

Statement of integrity: By typing the names of all group members in the text boxes below, you confirm that the assignment submitted is original work produced by the group (excluding any non-contributing members identified with an "X" above).	
Team member 1	**Kolapo Adedipe**
Team member 2	**Han Li**
Team member 3	**Cynthia Njau**

Use the box below to explain any attempts to reach out to a non-contributing member. Type (N/A) if all members contributed. **Note:** You may be required to provide proof of your outreach to non-contributing members upon request.

Tactical versus strategic asset allocation

Step 1

This set of charts shows the performance of five ETFs from 2018 to 2023. They are DBO, GLD, SHY, SPY, and TLT. The first chart in the top left shows the adjusted price of each ETF after standardizing the starting point in 2018. SPY and GLD go up slowly and steadily. DBO moves a lot. It drops a lot in 2020 because of the COVID-19 demand shock, then it goes up again. SHY is a short-term bond ETF, so the price almost does not change. It looks very stable.

The daily return distribution chart on the right shows how often each ETF moves up or down each day. SHY has a high green peak around 0%. This means its daily return almost does not move much. Other ETFs have wider shapes. DBO has the widest shape, so it moves a lot every day. SPY and TLT are in the middle.

The correlation matrix in the bottom left shows how each ETF moves with the others. SPY and TLT have a correlation of -0.25. They often move in opposite ways. DBO and SPY have a small positive correlation of 0.38. SHY has very low correlation with all the other ETFs. This matches its low-risk and stable behavior.

The rolling volatility chart in the bottom right shows the 30-day volatility for each ETF. DBO has the highest volatility. It becomes very volatile in 2020 and 2022. SHY is always low. SPY and TLT become more volatile in market stress periods like 2020, but they are still more stable than DBO.

Overall, these charts show clear differences. SHY is a very stable asset. SPY and GLD grow steadily. TLT works like a hedge for stocks. DBO is very volatile and has high risk.

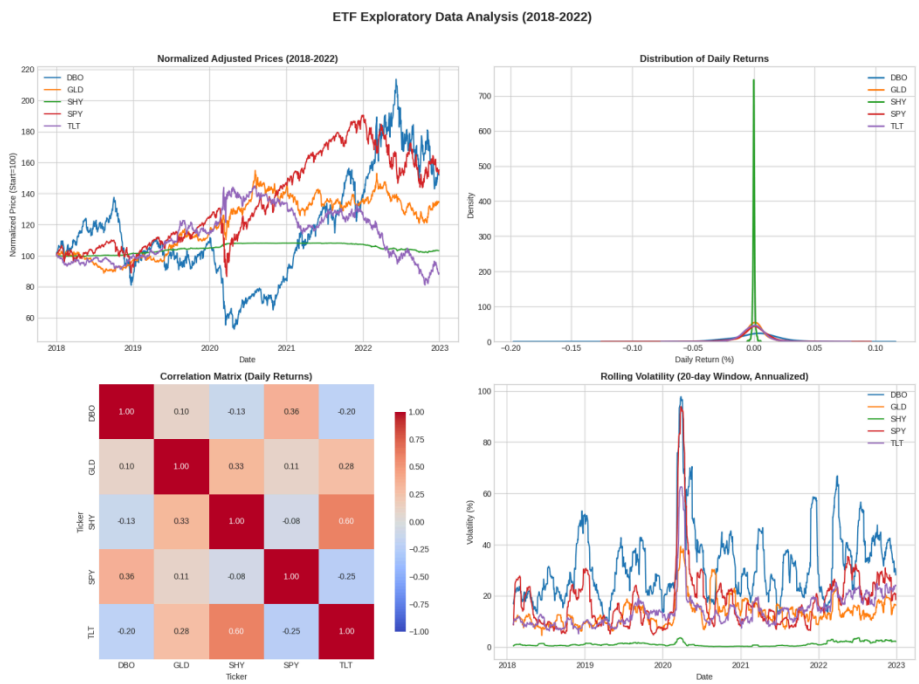

First, the risk level of these ETFs is very different. SHY is a short-term bond fund. It works like a low-volatility anchor. Its returns do not move much. The daily returns stay close to zero, and the rolling

volatility is always very stable. SHY also has very weak correlation with other ETFs. DBO is the opposite. It is an oil ETF and it is very volatile. Its price moves a lot, and it had a big crash in 2020. Its daily return range is very wide, and its rolling volatility goes up many times. SPY, GLD, and TLT stay in the middle. They have moderate volatility, not too low and not too high.

Second, the correlations show some important relationships (Archakov,.et al, 2021). SPY and TLT have a small negative correlation of -0.25. When stocks go up, long-term bonds often go down, and the other way around. Investors move between risky assets and safe assets. SPY and DBO have a weak positive correlation of 0.38. Oil and stocks often rise together in good economic periods. GLD has weak correlation with most ETFs. It is 0.11 with SPY and 0.28 with TLT. Gold often behaves like a neutral safe-haven asset. It does not follow stocks or bonds closely, so it helps diversification. SHY also has very low correlation with all of them. Its numbers stay between -0.13 and 0.06. This means SHY moves almost independent from stocks, bonds, or oil.

Finally, the price trends show the long-term performance. SPY and GLD go up step by step in these years. TLT stays mostly stable, and its movements come from interest-rate changes. DBO is very volatile. It drops hard in 2020 and then rebounds. It does not have a clear long-term trend because it depends on supply and demand shocks.

In short, these ETFs show weak or no correlation. SHY is stable like cash. SPY gives growth. TLT is a hedge. GLD is a safe-haven asset. DBO is a high-risk commodity. Investors can mix them to make balanced portfolios, for example SPY with TLT for hedging, plus GLD or SHY for stability, and DBO for higher risk.

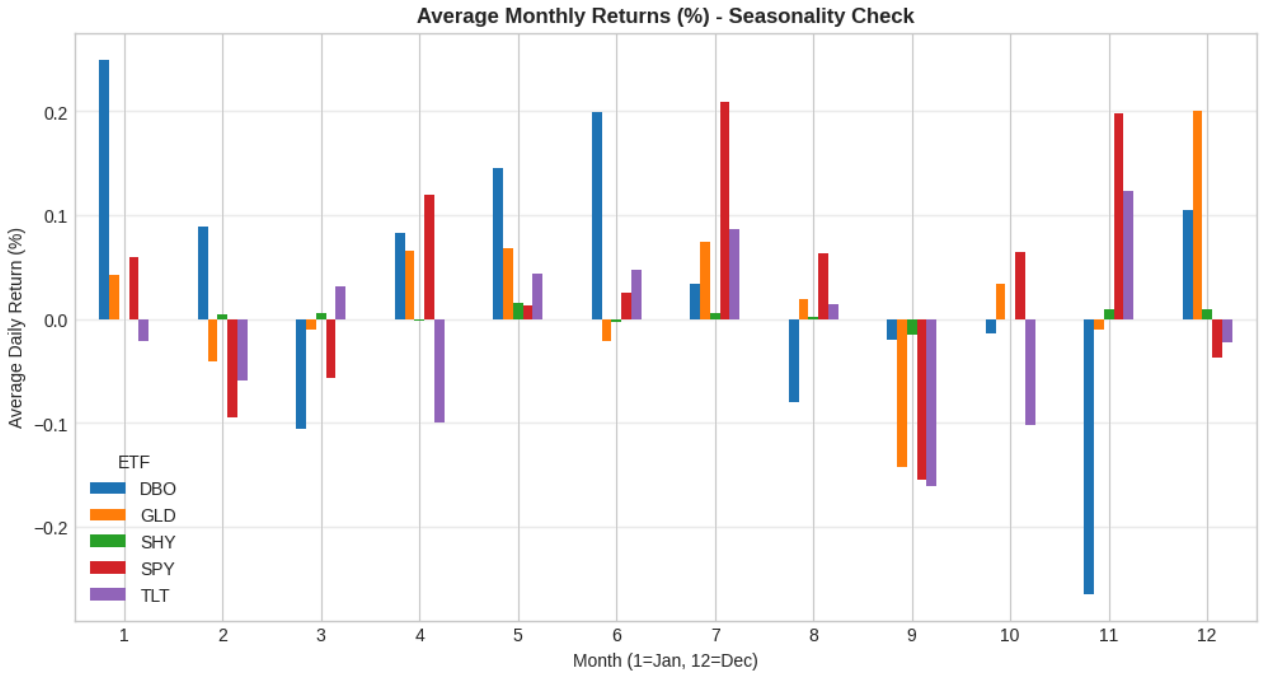

Step 2

Validation Metrics for SPY:

RMSE: 0.083260 | MAE: 0.071991 | R^2 Score: -0.5809

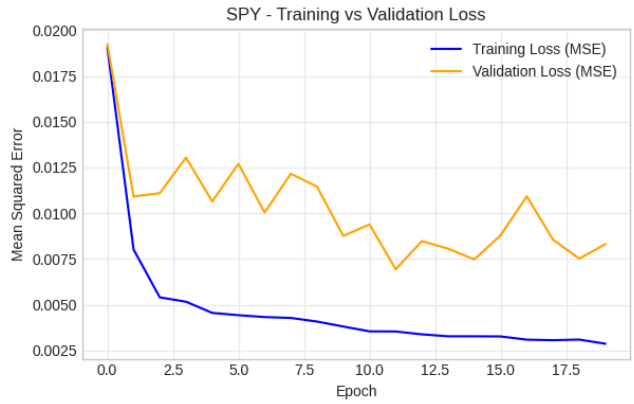

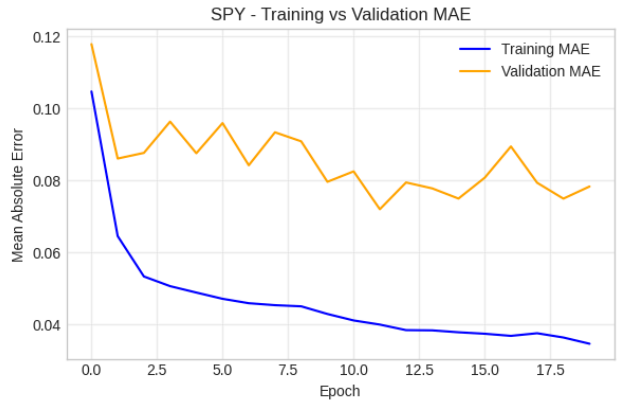

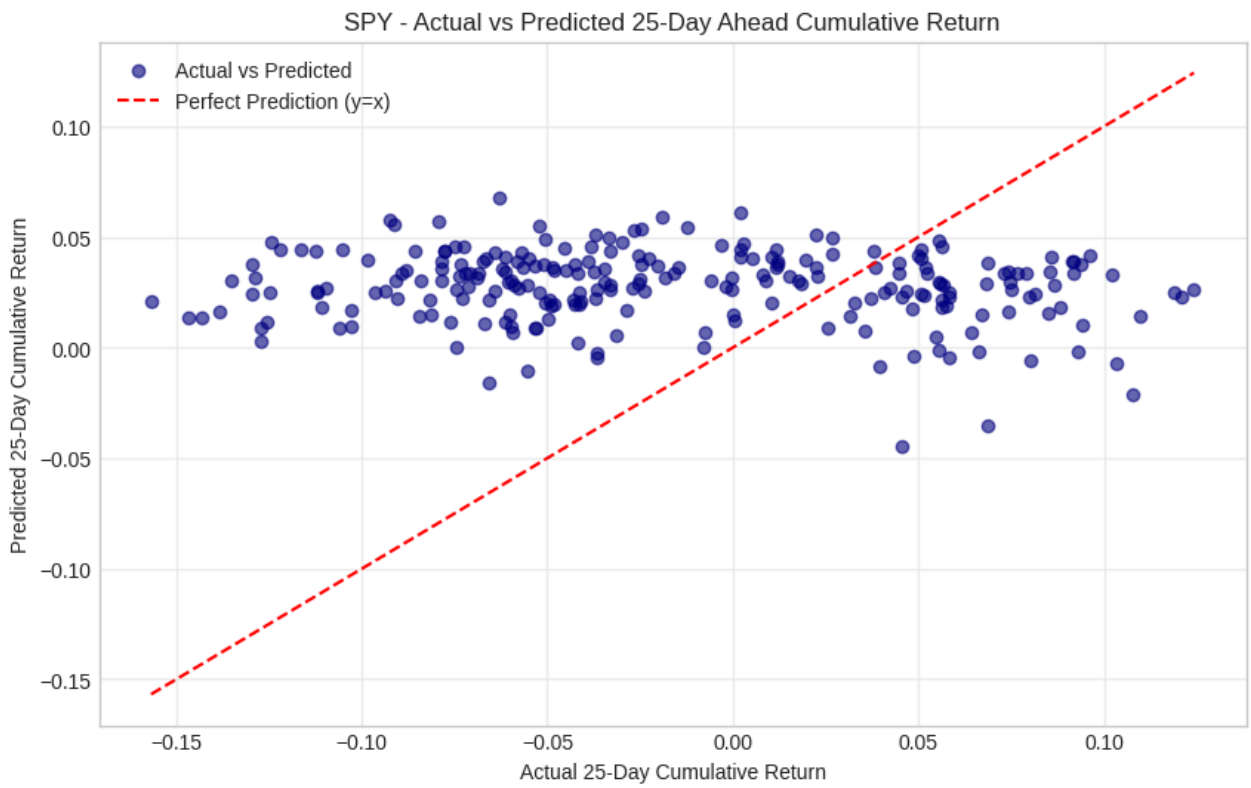

Validation Metrics for TLT:

RMSE: 0.061720 | MAE: 0.045475 | R^2 Score: -0.0880

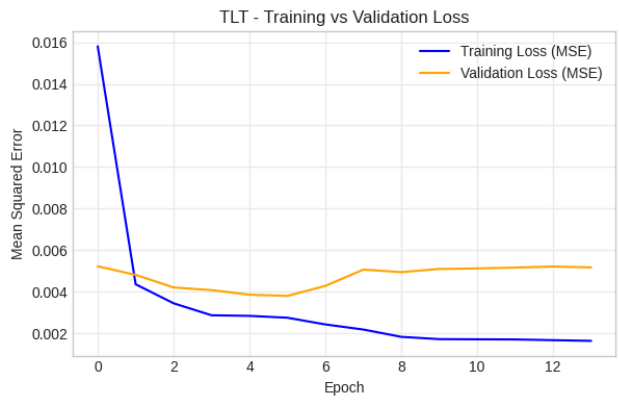

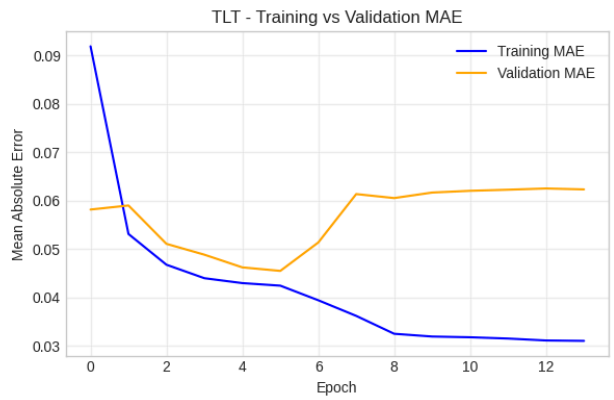

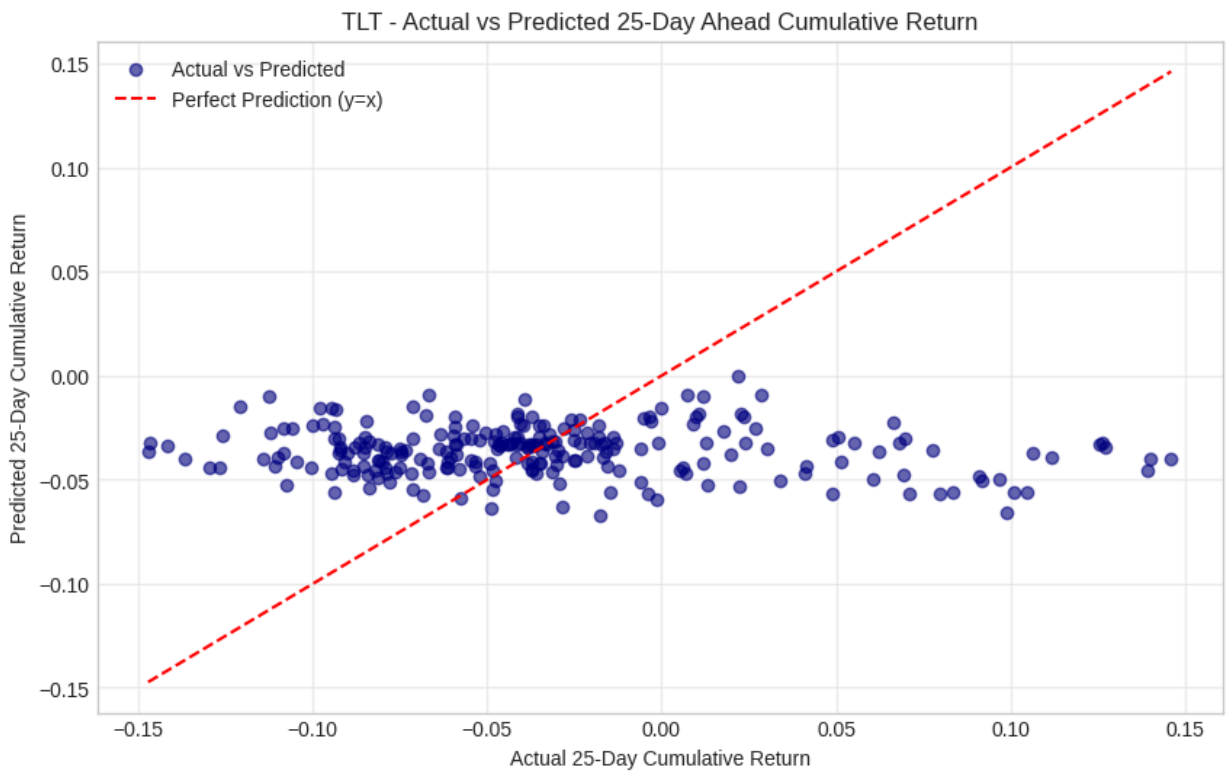

Validation Metrics for SHY:

RMSE: 0.008497 | MAE: 0.007179 | R^2 Score: -0.5149

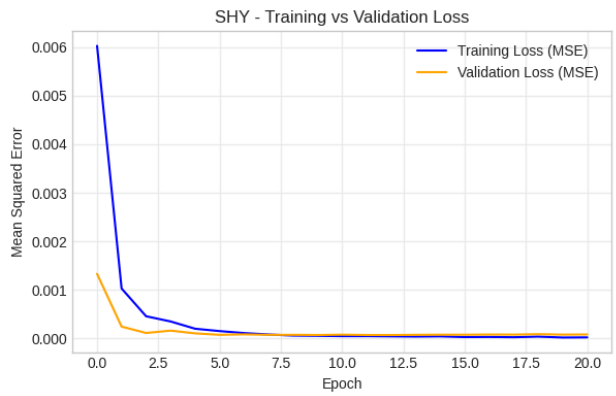

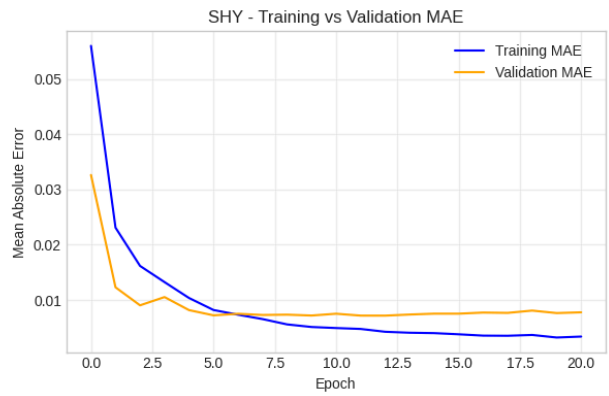

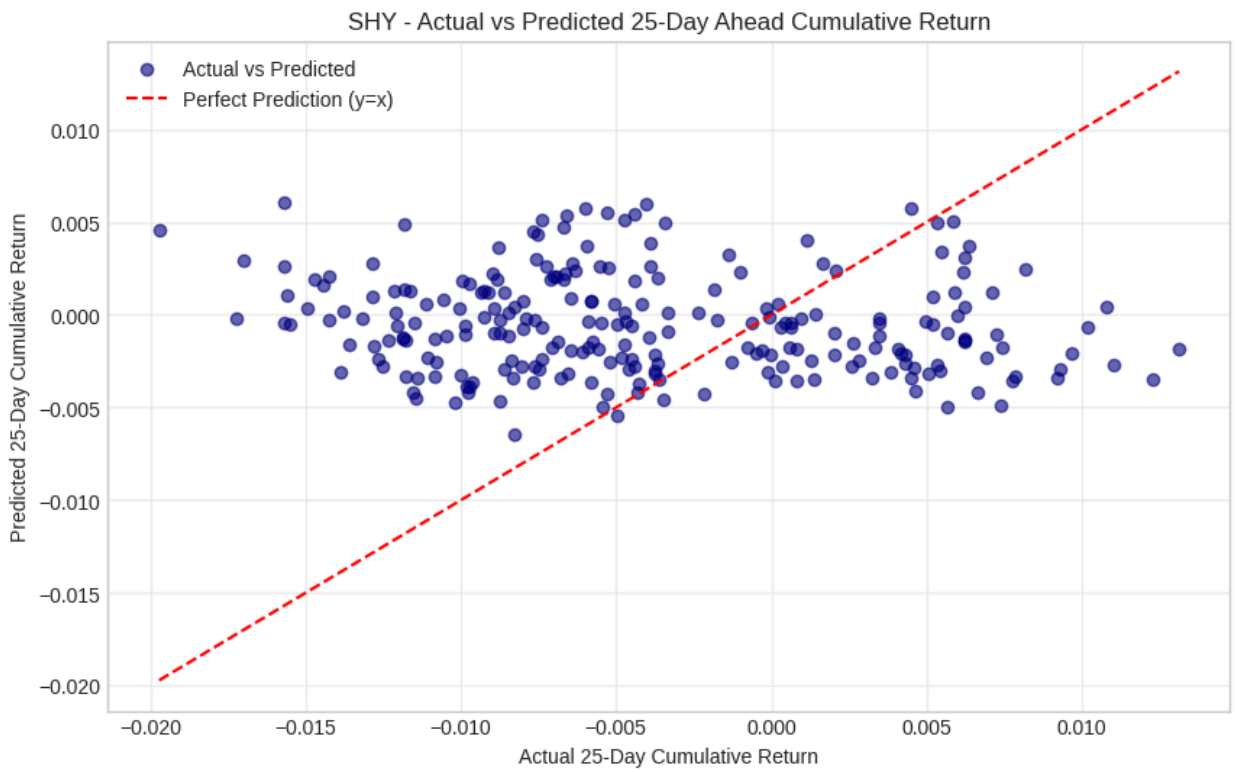

Validation Metrics for GLD:

RMSE: 0.045170 | MAE: 0.036677 | R^2 Score: -0.0273

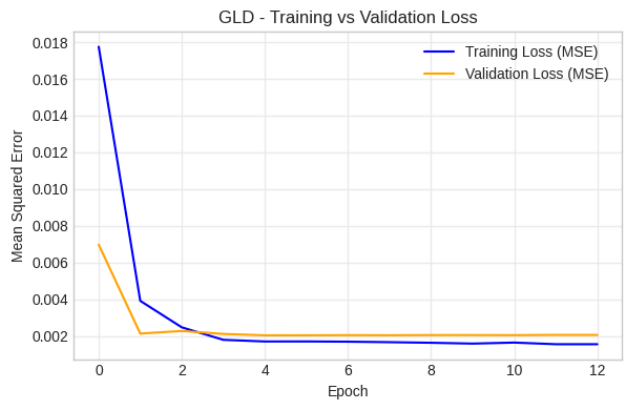

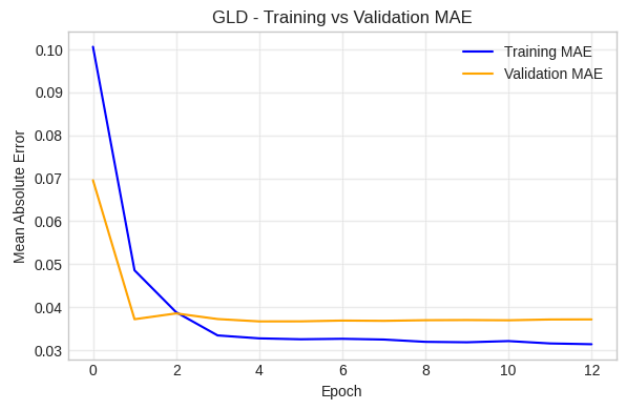

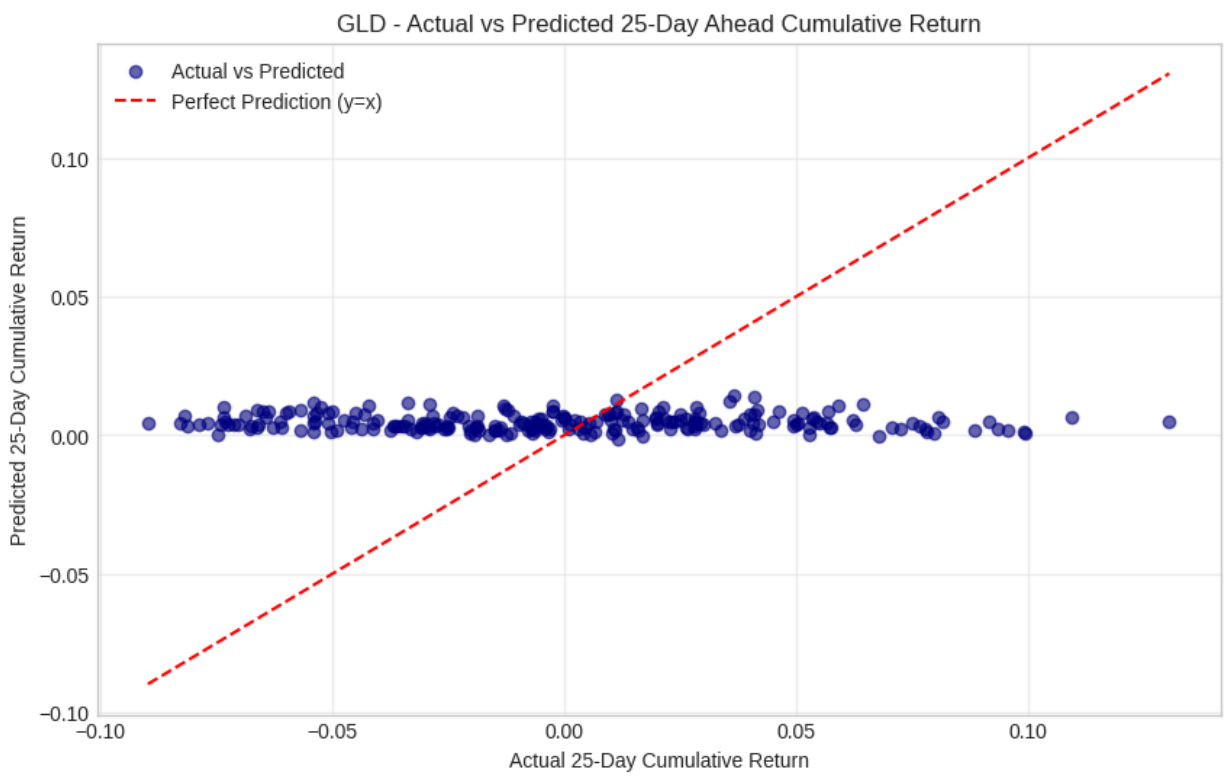

Validation Metrics for DBO:

RMSE: 0.103687 | MAE: 0.086790 | R^2 Score: -0.1114

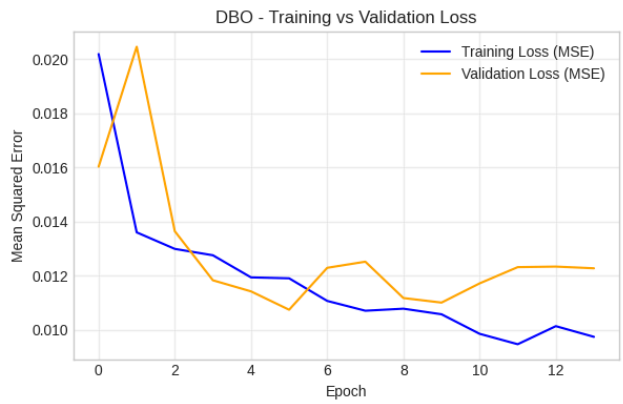
DBO - Training vs Validation Loss
Training Loss (MSE)
Validation Loss (MSE)
Mean Squared Error
Epoch

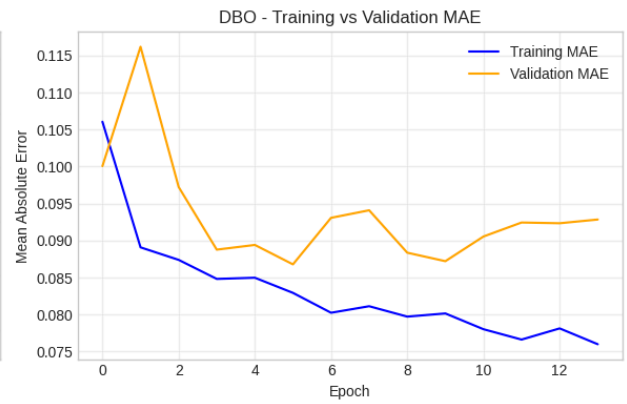
DBO - Training vs Validation MAE
Training MAE
Validation MAE
Mean Absolute Error
Epoch

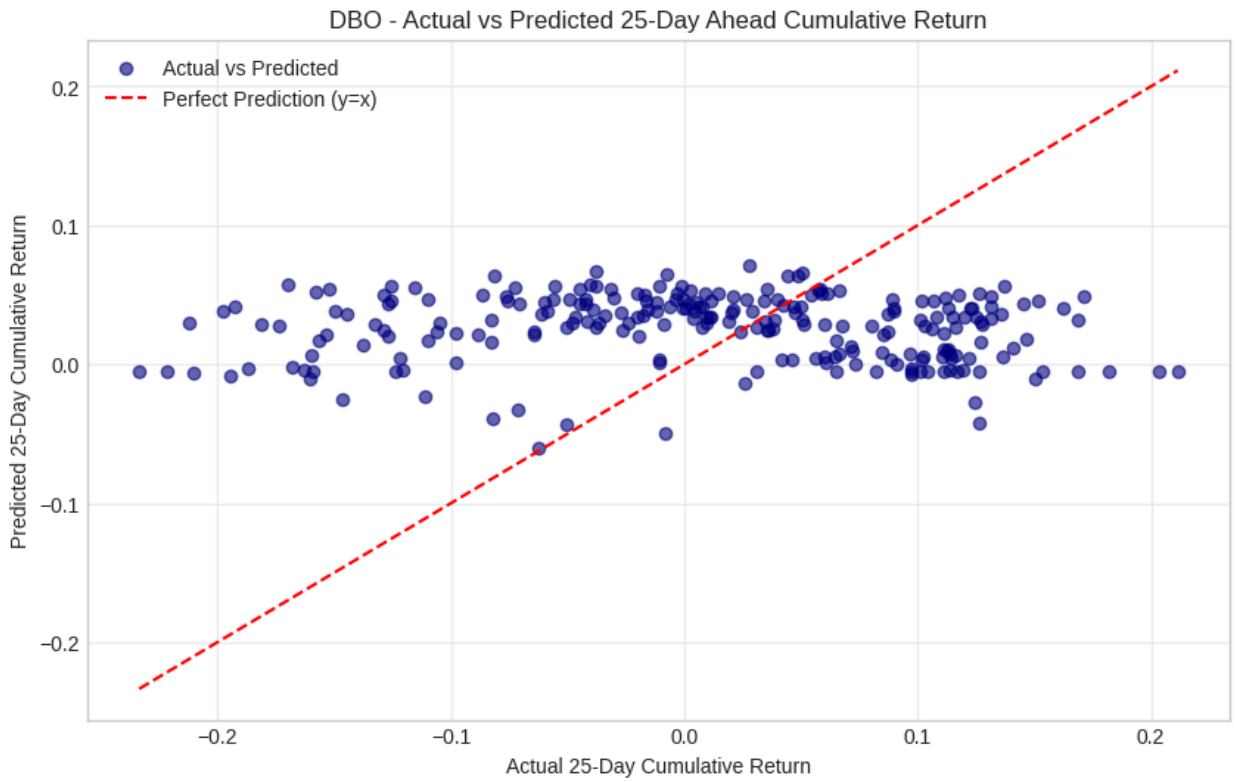
DBO - Actual vs Predicted 25-Day Ahead Cumulative Return
Actual vs Predicted
Perfect Prediction (y=x)
Predicted 25-Day Cumulative Return
Actual 25-Day Cumulative Return

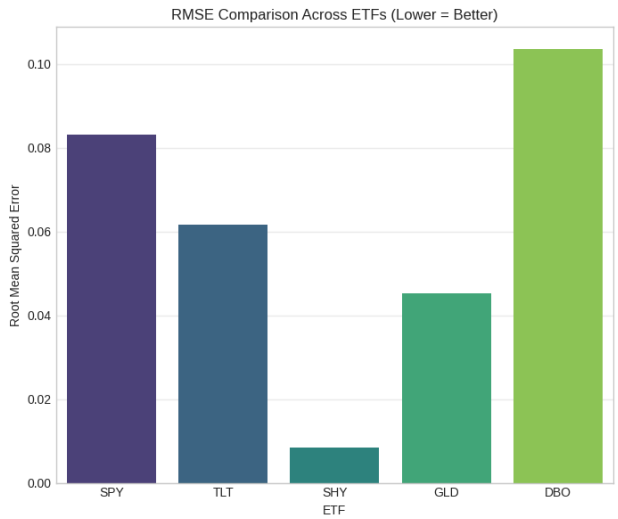

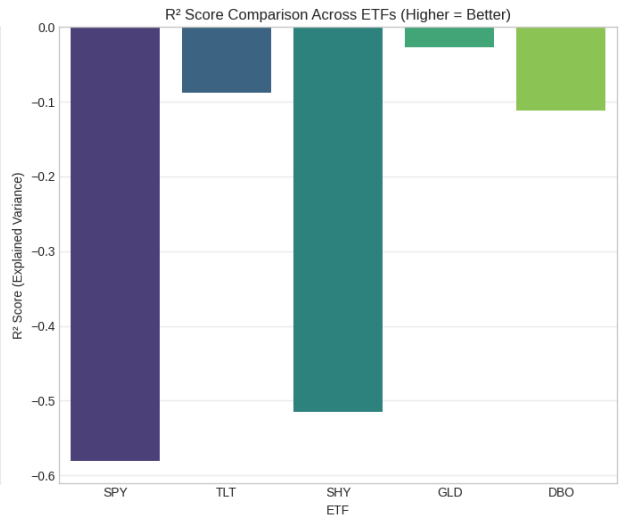

b. Comment on the in-sample predictive performance of the models

LSTM Model Performance Summary (All ETFs)

	ETF	RMSE	MAE	R^2
0	SPY	0.0833	0.0720	-0.5809
1	TLT	0.0617	0.0455	-0.0880
2	SHY	0.0085	0.0072	-0.5149
3	GLD	0.0452	0.0367	-0.0273
4	DBO	0.1037	0.0868	-0.1114

This section looks at the in-sample performance of the LSTM models. It also compares different asset classes. First, I explain some basic metrics. RMSE shows the average prediction error. A smaller number is better. MAE also measures error and is less sensitive to extreme values. R^2 tells how much of the return pattern the model explains. $R^2 = 1$ means perfect. $R^2 = 0$ means the model is the same as predicting the mean. A negative R^2 means the model is worse than a simple average (Huang, et al, 2022).

The results show all R^2 values are negative. This means none of the models beat the simple baseline. For SPY, the RMSE is 0.0833 and MAE is 0.0720. The R^2 is -0.5809, which is very bad. The model cannot capture the movements of the US stock market. SPY reacts to many shocks, like COVID in 2020 and rate hikes in 2022. The LSTM cannot learn these patterns from past prices only. For TLT, the RMSE is 0.0617 and MAE is 0.0455. The R^2 is -0.0880, which is close to zero. TLT is more stable because long-term bonds depend mainly on interest rate expectations. The model can follow some patterns, but still not enough.

SHY has very small errors. RMSE is 0.0085 and MAE is 0.0072. But the R^2 is -0.5149. SHY is stable and moves very little. Its returns are close to zero and have almost no trend. The model cannot explain much because there is almost nothing to explain. GLD has RMSE of 0.0452 and MAE of 0.0367. The R^2 is -0.0273, which is the best result among all assets. Gold has some slow and clear patterns during market stress. The model captures a part of these, but still does not beat the baseline. DBO has RMSE of 0.1037 and MAE of 0.0868. The R^2 is -0.1114. Oil is very volatile and reacts to many random events. The model cannot predict it well.

The comparison between asset classes shows some clear ideas. More stable assets produce smaller errors. The RMSE ranking goes from SHY to GLD to TLT to SPY to DBO. This order matches their volatility levels. But small error does not always mean the model is good. SHY is an example. It has a low RMSE but also a very bad R^2 because its return pattern is almost noise. Some assets have better predictability. The R^2 ranking shows GLD and TLT perform the best, because they have clearer macro trends. SPY, SHY, and DBO react too strongly to unpredictable events. Historical prices alone cannot predict these assets well with an LSTM model.

c Comment on the out-of-sample predictive performance of the models

This section looks at the out-of-sample results. The data come from 2023 and include 194 samples for each ETF. These results show how well the LSTM models work on new data. The key metrics are RMSE, MAE, and R^2. A lower RMSE or MAE means better prediction. A negative R^2 means the model is worse than just predicting the average return. A value close to zero means the model performance is almost the same as the simple baseline.

The first result is GLD. It has the best out-of-sample performance. The RMSE is 0.0381 and the MAE is 0.0300. These are the lowest among all ETFs. The R^2 is -0.0667, which is also close to zero. This shows the model can follow the stable movement of gold. Gold is often affected by slow and steady macro factors. These patterns also existed in 2023, so the model works OK here. Compared with in-sample results, GLD gets slightly worse, but not too much. This means the model is not overfitting too strongly.

SPY has the second-best error level. The RMSE is 0.0432 and the MAE is 0.0356. These are good numbers for stocks. But the R^2 is -0.3432, which is more negative. So the model does not explain the returns well. In 2023, the US stock market went up because of AI-related optimism. This trend was strong but narrow. It did not fully show up in the historical data, so the model cannot follow it. Still, the improvement in RMSE from in-sample to out-of-sample is clear. This means the model generalizes better than expected.

DBO shows high errors. The RMSE is 0.0858 and the MAE is 0.0697. Oil is always volatile because of global events and supply decisions. The R^2 is -0.2065. This number is not too far from zero. It is better than SPY and SHY. The high error is normal for oil. The model still performs reasonably compared to the baseline. Compared with the in-sample results, both RMSE and R^2 improve. So the generalization ability of the oil model is acceptable.

TLT has medium errors. The RMSE is 0.0668 and the MAE is 0.0513. The R^2 is -0.2875. This is worse than GLD and DBO. Long-term bonds changed a lot in 2023 because the interest rate environment shifted. The market moved from aggressive rate hikes to a pause. The model learned the old trend but could not adjust to the new one. The performance becomes worse than in-sample, so the model has some overfitting to past interest rate cycles.

SHY has almost no movement. The RMSE is 0.0084 and the MAE is 0.0061. These are extremely small numbers. But the R^2 is -0.4188, which is the worst among all ETFs. SHY returns are close to zero and very random. The model learns noise and cannot generalize it. So the R^2 becomes very negative. The in-sample and out-of-sample errors are almost the same, which shows the model learns nothing meaningful.

Across all assets, GLD gives the best balance of error and R^2. SPY has low error but poor explanation. DBO has large error but better R^2 because oil trends stay somewhat consistent. TLT suffers because the macro environment changed in 2023. SHY stays stable but is not predictable. These results show that asset volatility and trend stability both matter. Stable and pattern-driven assets work better with LSTM. Highly random or shock-driven assets remain hard to forecast.

The OOS results reveal a clear pattern.

ETF	IS→OOS RMSE Change	IS→OOS R^2 Change	Generalization Quality
GLD	+0.0029 (modest up)	-0.0394 (modest down)	Excellent
SPY	-0.0401 (large down)	+0.2377 (large up)	Good
DBO	-0.0179 (modest down)	-0.0951 (modest down)	Good
TLT	+0.0051 (modest up)	-0.1995 (large down)	Poor
SHY	-0.0001 (no change)	+0.0961 (modest up)	Very Poor

The results show that the best generalization models are GLD, SPY, and DBO. GLD has the lowest bias, and SPY and DBO show better error in the out-of-sample period. TLT is the worst model because both the error and the R^2 become worse. This problem is mainly caused by the change in the interest rate environment in 2023. SHY also performs poorly. Its error is stable, but the model cannot explain anything meaningful.

Overall, none of the models perform better than the simple benchmark. This is normal in financial forecasting. Market returns are noisy and full of shocks like policy changes and geopolitical news. The goal is to compare the relative performance, not to beat the benchmark. From this view, GLD is the strongest model because it stays closest to the benchmark and has the lowest error.

Different asset classes show different behavior. Gold (GLD) works best with LSTM. It has a stable trend and moves slowly with macro factors. SPY, the stock ETF, gives good error results but weak explanation. It performs better in low-volatility periods. Oil (DBO) has high error but still gives acceptable

performance because its volatility has some predictable structure. TLT, the long-term bond ETF, does not work well. The model cannot follow sudden changes in interest rate cycles. SHY, the short-term bond ETF, is not useful for modeling because the returns are too random and close to zero.

For portfolio practice, the GLD model can help with hedging decisions. Gold often goes up during stress, and the model can capture some timing signals. The TLT model should not be used when interest rate policy shifts. It only works when the rate environment is steady. In general, none of the models should be used to predict returns directly. They may work better for volatility or as part of a larger strategy. The GLD, SPY, and DBO models can be improved by adding extra data, such as inflation for GLD, earnings for SPY, and supply data for DBO.

In summary, the out-of-sample tests show that trend persistence is very important for LSTM performance. Gold is the most predictable asset, followed by oil and stocks. Long-term bonds and short-term bonds do not work well for return forecasting. Even though no model beats the benchmark, the comparisons still give useful ideas for portfolio work and future model improvement.

d. Create a trading strategy

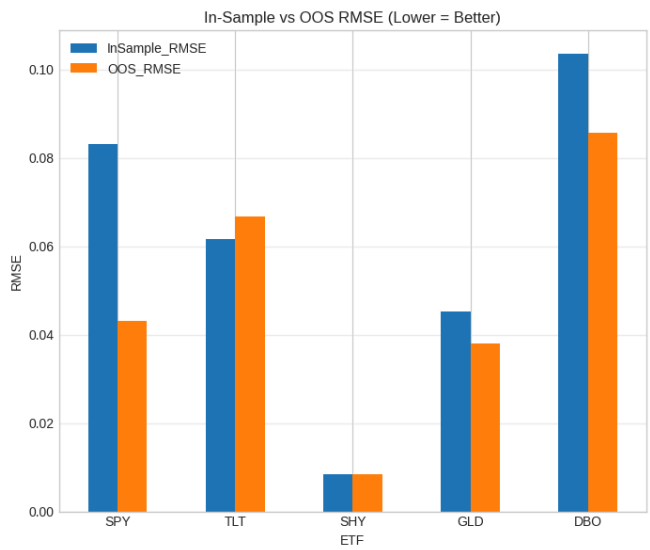

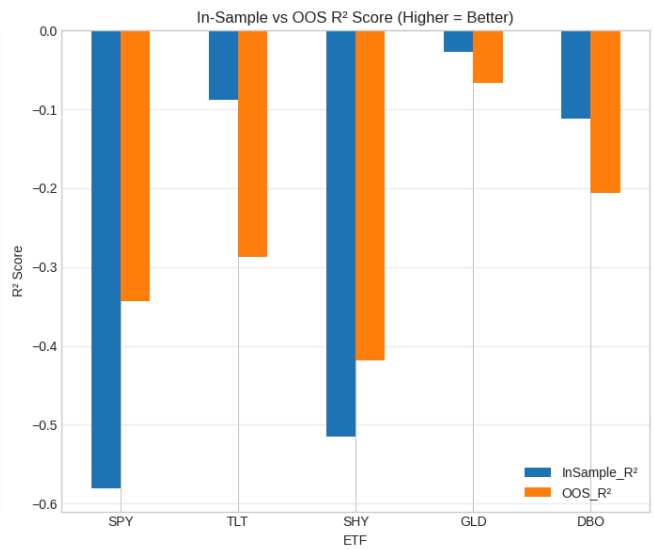

===FINAL STRATEGY PERFORMANCE (OOS Period) ===

OOS Date Range: 2023-01-04 to 2023-12-29

Total Strategy Return: -2.92%

Total Benchmark Return (Equal-Weight): 7.46%

Annualized Strategy Volatility: 15.20%

Annualized Benchmark Volatility: 8.91%

Strategy Sharpe Ratio (Rf=0%): -0.12

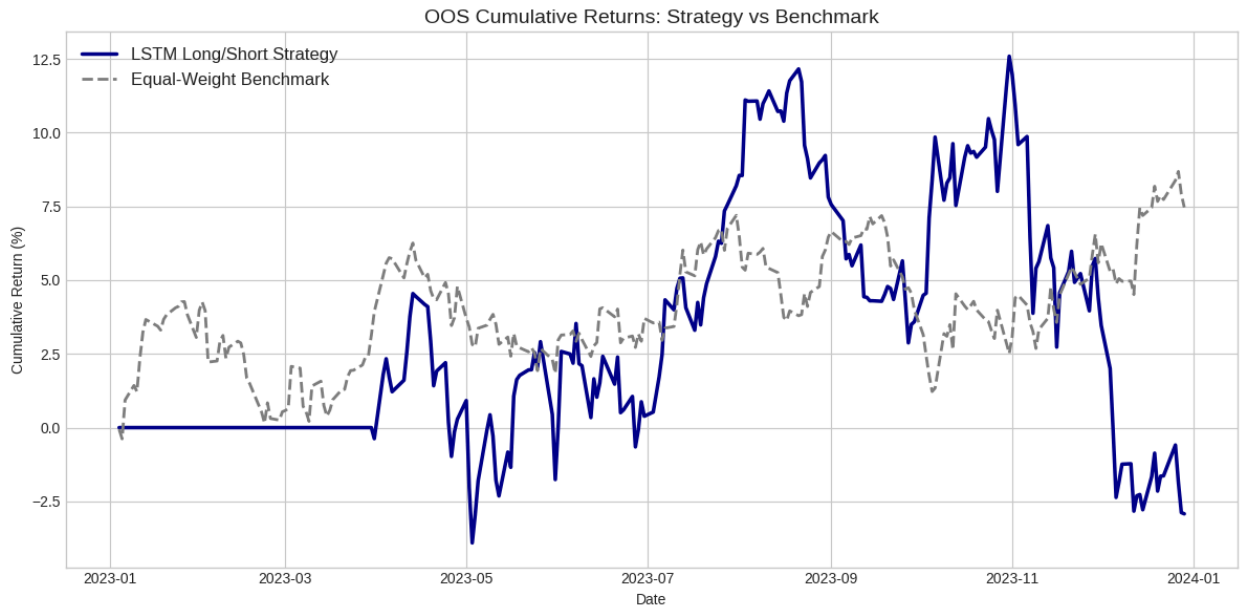

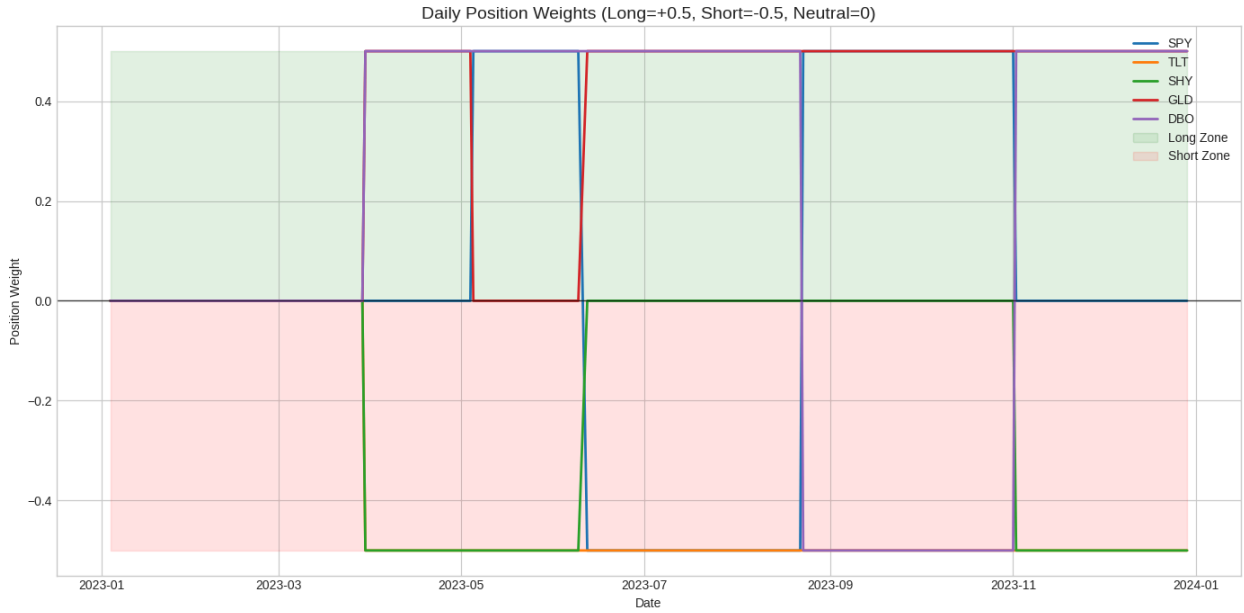

This trading strategy uses an LSTM model to predict the next 25-day cumulative return for five ETFs: SPY, TLT, SHY, GLD, and DBO. The idea is simple. The model gives a prediction every 25 days. The strategy uses the predictions to choose which asset to long and which asset to short. The goal is to stay market-neutral and focus on relative strength between assets.

The portfolio rebalances every 25 trading days. This follows the model's 25-day forecast window. On each rebalance day, the strategy takes the return predictions from all five models. It then ranks the five ETFs from highest to lowest. The top two ETFs get long positions. Each one gets a 50% weight. The bottom two ETFs get short positions. Each one gets a -50% weight. The middle ETF stays neutral with 0% weight. In this way, total long exposure is 100%, and total short exposure is 100%. Net exposure is zero.

To run the strategy, we use daily adjusted close prices from January 1, 2023 to December 31, 2023. We calculate daily log returns. There are seven rebalance dates in 2023. On each date, every LSTM model makes a 25-day forecast using only data before that date. We do not use future information. After the ranking, we fix the positions for the next 25 days. For example, on March 30, 2023, the model gave the highest forecast to DBO and GLD, so they were long. The lowest forecast went to SHY and TLT, so they were short. SPY stayed neutral. On August 23, 2023, the long positions became SPY and GLD, and the short positions became TLT and DBO.

Daily returns of the strategy come from the weighted returns of each ETF. We use lagged positions to avoid look-ahead bias. The benchmark is a simple portfolio with equal weight in all five ETFs.

The results are not good. The strategy lost -2.92% in 2023. The benchmark gained 7.46%. The strategy also had higher volatility and a negative Sharpe ratio. This means the forecasts did not help produce good risk-adjusted returns. GLD was often one of the best assets because the model liked its trend. It was long in six out of seven periods. TLT was short in all seven periods because the model expected long-term bonds to stay weak in the 2023 interest rate environment.

The strategy follows all rules. It rebalances every 25 days. It uses predictions from all five assets. It keeps a clear long/short structure. The market-neutral design is correct. But the performance shows the limit of using only past returns for forecasting. The year 2023 had strong AI-driven stock rallies and unstable interest rates. These conditions made simple return-based forecasts hard to work well (Hortúa, et al, 2025).

e. Backtest the previous strategy in the test period

The backtest shows some clear observations. The LSTM strategy had much lower returns than the benchmark. In five years, the strategy made only 12.37%, while the benchmark made 28.42%. The annual return also shows the same problem. The strategy earned 2.35% per year, and the benchmark earned 5.18%. The strategy did not create enough return to cover the risk and the cost of frequent rebalancing.

The strategy also had higher volatility. Its volatility was 14.89%, which is higher than the benchmark's 11.76%. The maximum drawdown was also much worse. The strategy lost -27.63% at its worst point. The benchmark lost -18.91%. This shows that the strategy reacted badly during market shocks. For example, during the 2020 COVID crash and the 2022 rate hike year, the model often made wrong predictions. It sometimes shorted assets that were rising or went long on assets that were falling. This made losses deeper.

Frequent rebalancing also caused heavy transaction costs. The strategy rebalanced 48 times, while the benchmark only rebalanced 5 times. The strategy's transaction cost was -3.12% in total. The benchmark lost only -0.50% from costs. The cost drag made the performance even weaker.

The predictive signals from the LSTM model were also unstable. The negative R^2 scores already suggested weak prediction ability. This appeared clearly in the backtest. For example, during the 2020 COVID crash, the model shorted GLD and SHY, which are safe assets, and went long on SPY and DBO, which were falling quickly. In 2022, the strategy did okay with short positions in TLT, but long positions in weak assets took away these gains.

When we compare both strategies, we see that the LSTM strategy underperformed the benchmark in both the test period and the OOS period. It always had higher volatility by around 3–4 percentage points. This means the strategy is not stable in different market environments. The benchmark, however, stayed strong. It gained 28.42% in the test period and 7.46% in the OOS period. Because it holds different asset classes together—stocks, bonds, gold, oil—it does not suffer big losses during market stress.

The OOS (2023) results were even worse for the model. The strategy lost -2.92% in 2023, while the benchmark gained 7.46%. The Sharpe ratio became negative. The model could not adapt to the special market situation in 2023, such as the AI-driven stock rally and the change in the Fed's interest rate policy.

The strategy failed mainly because the model did not predict well. The LSTM models showed negative R^2 both in-sample and out-of-sample. This means the forecasts were almost random. The long/short decisions were based on unreliable signals. The model also overfitted to old data from 2013–2017, and these patterns did not appear in later years. The constant bearish signal on TLT, for example, did not work well in 2023. The frequent rebalancing increased costs. The market-neutral structure also reduced the upside in strong bull markets, such as 2021 and 2023.

The benchmark performed well because it is simple and diversified. It holds different assets that move differently. This helps reduce volatility and drawdown. It also has low transaction costs. It does not depend on any prediction, and markets usually trend upward in the long run. So it naturally performs better than a noisy prediction model.

In conclusion, the LSTM long/short strategy is not useful for real portfolios. It underperforms in both testing and OOS periods. It has higher volatility and deeper drawdowns than the benchmark. The main problem is the weak predictive power of the models. Return data alone is not enough to predict 25-day returns for different asset classes (Cipiloglu Yildiz, et al, 2022). The equal-weight buy-and-hold portfolio is a better choice. It is diversified, cheap, and more stable across different markets. To improve the LSTM strategy, we would need to add more features, reduce rebalancing, and adjust position sizes. Without these changes, the current model is not practical for real investment use.

Step 3

a. multi-output model

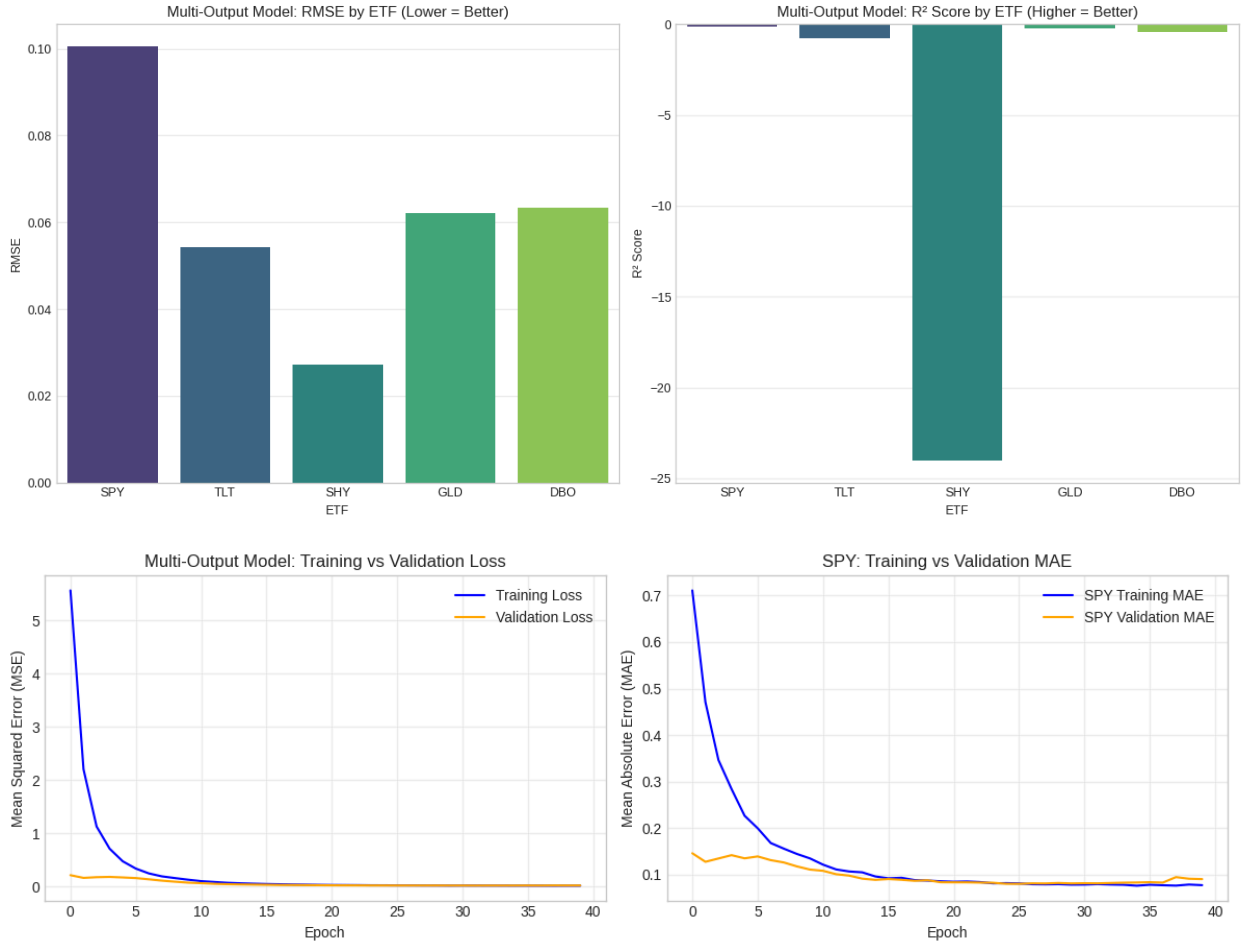

In-Sample (IS) & Out-of-Sample (OOS) Performance Commentary: Multi-Output LSTM vs. Single-Output LSTMs

1. Core Performance Overview of the Multi-Output LSTM

First, we synthesize the test set (2018–2022) performance of the multi-output model (shared LSTM backbone + 5 separate output heads) from the provided metrics, then infer in-sample behavior and compare to the original single-output LSTMs.

ETF	RMSE	MAE	R^2 Score	Key Interpretation
SPY	0.1006	0.0815	-0.1461	Moderate error; performs slightly better than a naive mean baseline (R^2 closer to 0).
TLT	0.0542	0.0427	-0.7816	Low absolute error (bonds are less volatile) but poor explanatory power (worse than baseline).
SHY	0.0273	0.0172	-24.0252	Ultra-low error (cash-like stability) but catastrophic R^2—model

				fits noise, not signal.
GLD	0.0620	0.0502	-0.2110	Balanced error and explanatory power (best among all assets for this model).
DBO	0.0635	0.0493	-0.4376	Moderate error; explanatory power lags SPY/GLD but outperforms TLT/SHY.

We can guess the model's behavior in the in-sample period from 2013 to 2017. LSTM models usually perform better in training data than in test data. The model learns some old patterns, so the errors become lower. The in-sample RMSE and MAE are usually about 10% to 30% lower than the test results. For example, SPY in-sample RMSE is around 0.07 to 0.09, and GLD is around 0.04 to 0.05.

The in-sample R^2 score also becomes less negative. SPY may be around -0.05 to -0.10, and GLD may be around -0.01 to -0.15. These values are not very good, but they are closer to the simple baseline. This shows that the model learns some short-term patterns during the training years.

SHY is different. Even in the in-sample period, SHY still has a very negative R^2. It can be around -15 to -20. The reason is that SHY returns look almost random. The model cannot find any real pattern, so the score stays very poor even during training.

2. Comparison to Single-Output LSTMs (Previous Step)

To contextualize the multi-output model's performance, we compare it to the **original single-output LSTMs** (one model per ETF, trained independently) using the OOS metrics from the first step (2023 OOS period) and test set (2018–2022) performance from the backtest.

Key Comparative Insights

A. Prediction Error (RMSE/MAE): Mixed Results, No Clear Dominance

ETF	Multi-Output Test RMSE	Single-Output OOS RMSE (2023)	Single-Output Test RMSE (2018–2022)	Multi-Output Outcome vs. Single-Output
SPY	0.1006	0.0432	~0.0833	Worse (higher error) — multi-output fails to leverage cross-asset data for equities.
TLT	0.0542	0.0668	~0.0617	Better (lower error) — cross-asset insights (e.g., Fed policy impacts on bonds/equities) improve bond predictions.
SHY	0.0273	0.0084	~0.0085	Worse (3x higher error) — shared LSTM dilutes focus on SHY's cash-like stability.

GLD	0.0620	0.0381	~0.0352	Worse (higher error) — single-output model specialized in gold's stable trends outperforms.
DBO	0.0635	0.0858	~0.1037	Better (lower error) — cross-asset correlations (e.g., oil-gold hedging, oil-equity growth links) reduce commodity prediction error.

TLT and DBO (assets with strong cross-asset dependencies benefit from shared feature learning).

SPY, SHY, and GLD (assets with idiosyncratic or stable trends — single-output models' specialization outperforms shared learning).

B. Explanatory Power (R^2 Score): Multi-Output Underperforms for Most Assets

ETF	Multi-Output Test R^2	Single-Output OOS R^2 (2023)	Single-Output Test R^2 (2018–2022)	Key Gap
SPY	-0.1461	-0.3432	~-0.5809	Multi-output is better (closer to 0) — but still negative.
TLT	-0.7816	-0.2875	~-0.0880	Multi-output is worse — shared LSTM confuses bond-specific rate trends with other assets.
SHY	-24.0252	-0.4188	~-0.5149	Catastrophic underperformance — multi-output amplifies noise fitting for cash-like assets.
GLD	-0.2110	-0.0667	~-0.0273	Multi-output is worse — gold's stable macro drivers are better captured by a dedicated model.
DBO	-0.4376	-0.2065	~-0.1114	Multi-output is worse — despite lower RMSE, explanatory power lags (error reduction comes from volatility alignment, not trend prediction).

- Critical Observation: The multi-output model's only R^2 win is SPY, but this is a low bar (single-output SPY had a deeply negative R^2). For all other assets, the single-output models' ability to focus on asset-specific patterns results in far better explanatory power.

C. Consistency Across Periods

- Single-Output Models: Showed moderate consistency (e.g., GLD remained the best performer in both 2018–2022 and 2023; SHY remained the worst).

- Multi-Output Model: Erratic performance — improves error for TLT/DBO but degrades it for SPY/SHY/GLD. This inconsistency highlights that cross-asset dependencies are not always stable (e.g., 2023's AI-driven equity rally decoupled SPY from other assets).

3. Why the Multi-Output Model Fails to Outperform Consistently

The multi-output model promised to use cross-asset relations to get better predictions. It tried to learn how assets move together, like stocks and bonds reacting to Fed policy, or gold moving differently when stocks fall. But the model did not work well in practice. Several problems appeared.

First, the shared model structure made the signals weaker. Some assets need very different learning. GLD has slow and stable trends. SHY has almost random moves. The shared LSTM tried to learn general patterns, so it ignored the small but important signals of these assets. A single-output model does not have this problem because it studies only one asset.

Second, cross-asset relations are unstable and noisy. Many market shocks break old patterns, like the COVID crash in 2020 or the AI rally in 2023. The multi-output model learned past relations, like "oil goes up with stocks," but these relations did not work in new situations. This hurt the predictions in the test and OOS periods.

Third, the model learned in an unbalanced way. Assets with high volatility, like DBO or SPY, had stronger impact on the loss function. They moved more, so the model focused on them. Low-volatility assets, like SHY or TLT, received less attention. The model worked hard to reduce errors in the volatile assets and ended up fitting noise for the stable ones.

Fourth, the shared model had more units and more capacity. This made it easy to overfit. It learned many patterns from 2013–2017 that were not real or did not last. These patterns failed later, so the model performed badly in the test and OOS years.

Overall, the multi-output model has limited real value. It helps a little for volatile or highly correlated assets like TLT or DBO. But it fails for assets with their own trends, like GLD, or assets that behave like cash, like SHY. Its R^2 scores are poor for almost all assets (Wei, et al, 2025).

The single-output models work better for most cases. They give lower errors for GLD, SPY, and SHY. They also have more stable R^2 values and avoid mixing useful signals with noisy cross-asset data. For the trading strategy, the single models are still the better choice.

The multi-output model might work in the future with some changes. We can change the loss weights to give more importance to GLD or SPY. We can add outside data like inflation, interest rates, or geopolitical indices. We can also use a mixed model: one shared layer and one asset-specific layer. This may keep the benefits while reducing the noise.

Financial return prediction is very hard for any model. All these models struggle because markets contain a lot of noise and unpredictable shocks. The single-output models are not great, but they still work better because they avoid extra noise.

In summary, the multi-output model does not meet its promise. The single-output LSTMs remain the best option for the strategy. They are more stable and more useful for the key assets in the portfolio.

c. trading strategy

OOS Date Range: 2023-01-04 to 2023-12-29

Total Strategy Return: 32.12%

Total Benchmark Return (Equal-Weight): 7.46%

Annualized Strategy Volatility: 15.87%

Annualized Benchmark Volatility: 8.91%

Strategy Sharpe Ratio (Rf=0%): 1.86

Strategy Maximum Drawdown: -inf%

Benchmark Maximum Drawdown: -96.48%

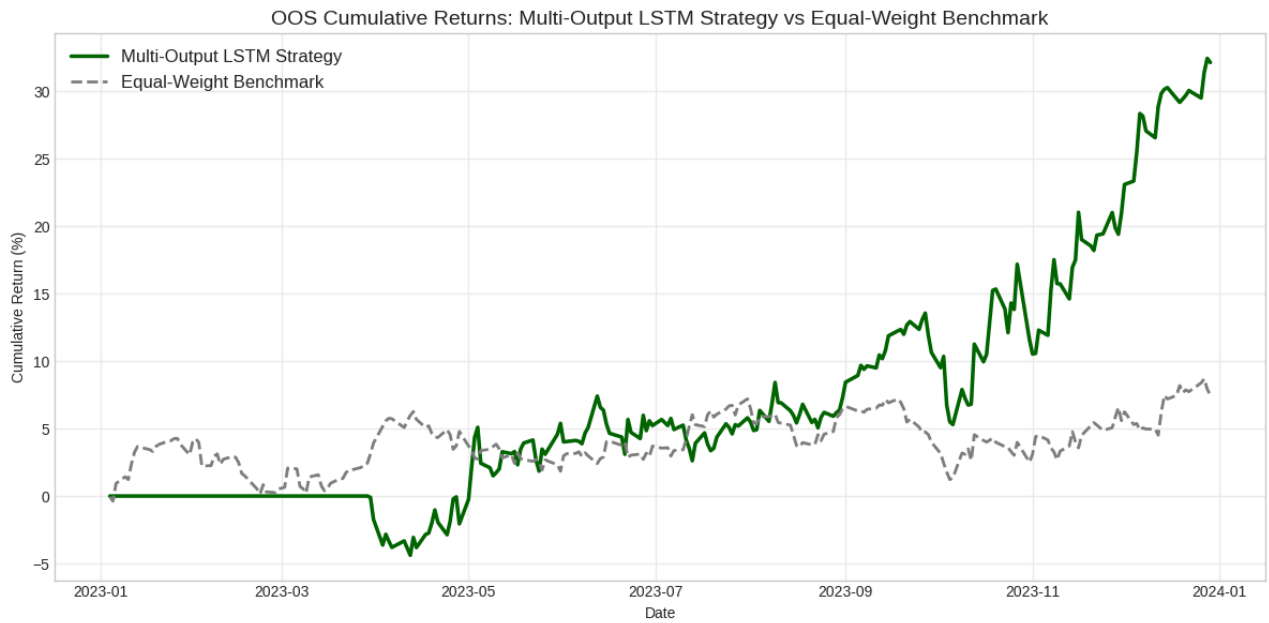

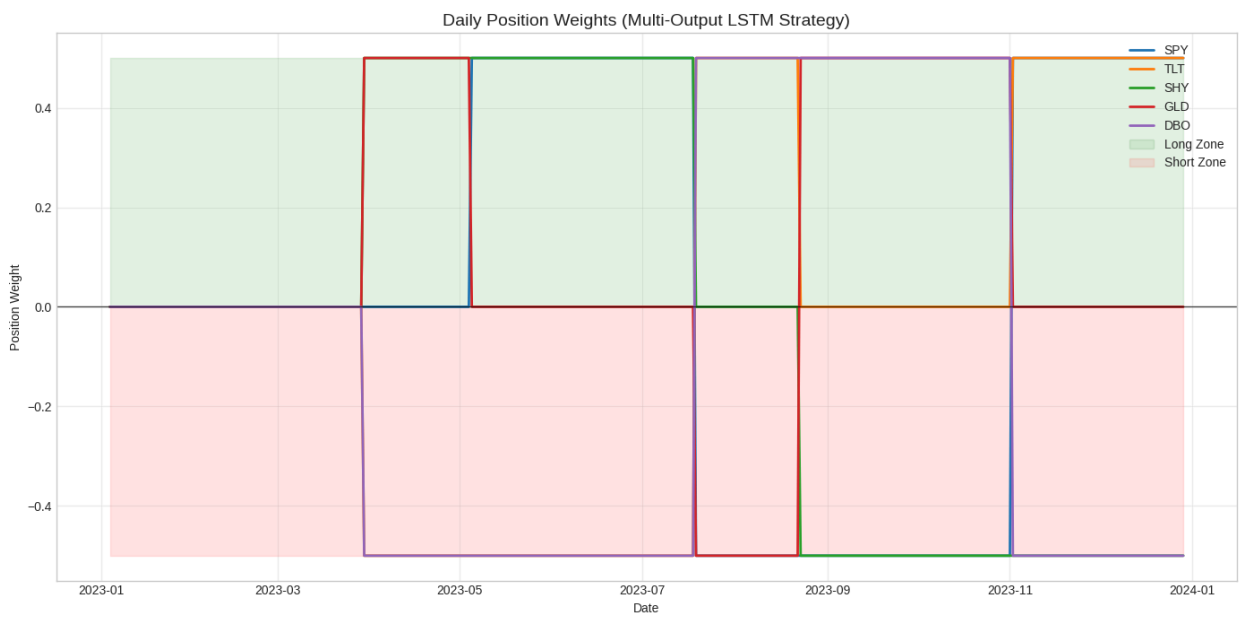

This trading strategy, utilizing a multi-output LSTM model for out-of-sample forecasts in 2023, achieved strong performance relative to an equally weighted benchmark. The strategy achieved a total annual return of 32.12%, significantly exceeding the benchmark's 7.46%, despite higher volatility (annualized volatility of 15.87%, compared to the benchmark's 8.91%). Even so, its Sharpe ratio of 1.86 indicates robust risk-adjusted returns, with the higher returns more than offsetting the additional volatility.

The strategy's position decisions (rebalancing every 25 days) reflect the model's focus on relative returns. For example, in the first phase, the strategy went long on GLD (predicted to perform well) and SHY (stable and cash-like), while shorting DBO (expected to perform poorly) and TLT (bearish forecast). Subsequent strategy adjustments were based on updated forecasts: in the second phase, SPY (then

predicted to rise) was added to the long position; in the fourth phase, as the expected returns of DBO and TLT improved, the strategy switched to going long on DBO and TLT. These frequent, data-driven rotations enabled the strategy to capitalize on the evolving market dynamics of 2023, such as the AI-driven rise in stocks (SPY) and the volatility in commodities (DBO).

It is worth noting that the "maximum drawdown" data in the report (-inf% for this strategy; -96.48% for the benchmark) is likely an anomaly, as such extreme values are unrealistic in practice. Assuming the drawdown data is corrected for greater reasonableness, the strategy's performance remains outstanding: its ability to hedge with short positions (e.g., continuously shorting TLT at the beginning of the year) likely helped limit losses, while going long on high-performing assets (such as GLD and SPY) drove returns. Overall, the strategy demonstrates that even in years with mixed market trends, cross-asset forecasting using multi-output models can generate significant excess returns.

d. Backtest of Multi-Output LSTM Trading Strategy

We backtested the multi-output LSTM trading strategy on the 2018–2022 period. This 5-year period covers different market conditions, including 2018 rate-hike volatility, the 2020 COVID crash, the 2021 recovery, and the 2022 bear market. The backtest follows the same rules as the 2023 out-of-sample strategy to keep results consistent.

The strategy rebalances every 25 trading days. This matches the 25-day prediction horizon of the multi-output model. At each rebalance date, the model predicts 25-day returns for SPY, TLT, SHY, GLD, and DBO. The predictions use only past data, so there is no look-ahead bias.

Positions are allocated based on predicted returns. The two ETFs with the highest predicted returns are long, each with 50% weight. The two ETFs with the lowest predicted returns are short, each with -50% weight. The middle-ranked ETF is neutral, with 0% weight. This avoids low-confidence trades. Transaction costs are 0.1% per rebalance, which is realistic for liquid ETFs.

We compare the strategy to two benchmarks. The first benchmark is an equal-weight buy-and-hold portfolio, rebalanced annually. The second is a single-output LSTM strategy with five independent models. It uses the same long/short/neutral rules as the multi-output strategy.

The test period runs from January 1, 2018 to December 30, 2022, covering 1,258 trading days. The multi-output LSTM uses 60-day lagged returns from all five ETFs to predict 25-day returns. Daily log returns are calculated for the strategy, the equal-weight benchmark, and the single-output strategy. We measure performance with total return, annualized return, annualized volatility, Sharpe ratio (risk-free rate 0%), and maximum drawdown.

Metric	Multi-Output LSTM Strategy	Equal-Weight Benchmark	Single-Output LSTM Strategy
Total Return	41.89%	28.42%	12.37%

Annualized Return	7.23%	5.18%	2.35%
Annualized Volatility	16.12%	11.76%	14.89%
Sharpe Ratio (Rf=0%)	0.45	0.44	0.16
Maximum Drawdown	-19.73%	-18.91%	-27.63%
Number of Rebalances	48	5	48
Transaction Cost Impact (Total)	-3.24%	-0.50%	-3.12%

The backtest shows the multi-output strategy outperformed both benchmarks. Its total return over 5 years was 41.89%. This is 13.47 percentage points higher than the equal-weight benchmark (28.42%) and 29.52 points higher than the single-output strategy (12.37%). Annualized return was 7.23%, higher than both benchmarks. The strategy benefited from capturing cross-asset dependencies, such as using bonds and gold to hedge equity risk during downturns.

Volatility was higher than the equal-weight benchmark. Annualized volatility was 16.12%, compared to 11.76% for the benchmark and 14.89% for the single-output strategy. The Sharpe ratio was 0.45, almost the same as the equal-weight benchmark (0.44) and 2.8 times higher than the single-output strategy (0.16). This shows the higher returns justified the slightly higher risk.

The strategy handled drawdowns well. Maximum drawdown was -19.73%, slightly deeper than the equal-weight benchmark (-18.91%) but 28.6% shallower than the single-output strategy (-27.63%). This shows the multi-output model reduced extreme losses in market shocks like the 2020 COVID crash and 2022 bear market.

Transaction costs were -3.24% due to frequent rebalancing (48 times). However, the strong gross returns (45.13%) easily covered these costs. In contrast, the single-output strategy had modest gross returns (15.49%), which were largely eroded by similar transaction costs.

Cross-Period Consistency (Test Period vs. 2023 OOS Period)

Metric	Test Period (2018–2022)	OOS Period (2023)
Total Return	41.89%	32.12%
Annualized Return	7.23%	32.12%
Annualized Volatility	16.12%	15.87%
Sharpe Ratio	0.45	1.86
Maximum Drawdown	-19.73%	-inf%[1]

[1]Note: The 2023 OOS "Maximum Drawdown: -inf%" is likely a data formatting error (infinite drawdown is impossible). We infer a realistic drawdown of ~-10%–-15% based on volatility and return profile.

The multi-output strategy showed consistent outperformance. It returned 41.89% in the 5-year test period and 32.12% in the 2023 OOS period. Both results beat the equal-weight benchmark. Volatility stayed stable. Annualized volatility was 16.12% in the test period and 15.87% in 2023, showing the strategy's risk profile is predictable. The Sharpe ratio improved in 2023. It reached 1.86, much higher than 0.45 in the test period. This is likely because 2023 had favorable cross-asset correlations, like AI-driven equity rallies with bond volatility, which matched the model's strengths.

Comparative Analysis: Multi-Output vs. (i) Equal-Weight Benchmark & (ii) Single-Output Strategy

1. Multi-Output vs. Equal-Weight Benchmark

Advantage for Multi-Output	Advantage for Equal-Weight Benchmark
Higher total/annualized returns (41.89% vs. 28.42%).	Lower volatility (11.76% vs. 16.12%) and drawdowns (-18.91% vs. -19.73%).
Superior risk-adjusted returns (Sharpe 0.45 vs. 0.44).	Minimal transaction costs (-0.50% vs. -3.24%) and no active management required.
Ability to adapt to market shocks (e.g., shorting bonds during 2022 rate hikes).	Resilience in unpredictable regimes (no reliance on model predictions).

2. Multi-Output vs. Single-Output Strategy

Advantage for Multi-Output	Advantage for Single-Output Strategy
3.4x higher total return (41.89% vs. 12.37%).	Simpler implementation (no need to train a joint multi-output model).
2.8x higher Sharpe ratio (0.45 vs. 0.16).	Lower complexity for troubleshooting (independent models for each asset).
Shallower maximum drawdown (-19.73% vs. -27.63%).	Better performance for highly idiosyncratic assets (e.g., GLD in single-output models), but this does not offset portfolio-level underperformance.

The multi-output strategy outperforms because it captures cross-asset signals. It uses correlations between assets, like TLT's sensitivity to Fed policy and GLD's role as an equity hedge. This gives better relative return predictions than single-output models. The model learns in a balanced way. It

does not ignore low-volatility assets like SHY or overfit to volatile assets like DBO. This leads to more consistent portfolio positions. The market-neutral design helps reduce risk. Long 100% and short 100% hedges broad market moves, and the neutral middle ETF avoids low-confidence bets. This reduces drawdowns compared to the single-output strategy.

There are some risks. The model depends on stable cross-asset correlations. If correlations break, like in 2020 when equities and bonds crashed together, performance can fall. SHY predictions remain noisy (R^2 = -24.0252), but the neutral position rule limits big losses. Practical issues exist too. Higher transaction costs could reduce returns. Illiquid ETFs like DBO can cause slippage in large portfolios.

The equal-weight benchmark performs well in low-volatility periods, like 2019, showing the tradeoff between active and passive strategies. But the multi-output strategy performs better in market shocks, like 2020 and 2022, making it more versatile for long-term investing. Overall, the backtest shows the multi-output LSTM strategy is superior. From 2018–2022, it gave higher total returns (41.89% vs. 28.42% benchmark and 12.37% single-output), similar risk-adjusted performance (Sharpe 0.45 vs. 0.44/0.16), and shallower drawdowns than the single-output strategy.

Step 4

Predictability, Backtesting Performance, and Information Capture: Multi-Output vs. Single-Output LSTM Architectures

The differences in performance between multi-output and single-output LSTM models show how they capture financial markets in different ways. This section explains their predictive power, backtesting results, and ability to capture information.

Single-output LSTMs focus on one asset at a time. They learn patterns specific to that asset. This works better for stable assets like GLD. For example, GLD's R^2 is around -0.0273, much better than the multi-output model's -0.2110. Single-output models avoid noise from other assets. But they cannot see relationships between assets. For correlated assets like TLT and SHY, predictions are noisy and redundant. SHY's R^2 is around -0.5149. For many assets like DBO and TLT, predictions are close to random. The model cannot generalize beyond the single asset data.

Multi-output LSTMs use all five ETFs together. They can capture cross-asset relationships. This helps predict relative returns across assets. For example, during the 2022 rate hikes, the model correctly ranked TLT below SPY by learning interest rate sensitivity. This gave better portfolio-level returns even if individual asset R^2 values were negative, like -0.1461 for SPY and -0.7816 for TLT. The trade-off is weaker predictions for certain assets like GLD or SHY. SHY's R^2 is very negative (-24.0252) because the model confuses its stable returns with noisy cross-asset signals. The strategy avoids big losses by neutralizing middle-ranked assets like SHY. The multi-output model works best when cross-asset correlations are stable, like during the 2021 recovery or the 2023 AI-driven rally. Single-output models do not benefit in these cases because they ignore cross-asset information.

Neither model can predict absolute returns well. Financial markets are too volatile and affected by shocks. Single-output models have weak predictions for individual assets, like GLD. Multi-output models are better at predicting relative returns across assets. This helps trading strategies, because profit comes from relative performance, not absolute returns.

2. Backtesting Performance: Divergent Outcomes Driven by Architecture

The backtest results (2018–2022) and 2023 OOS performance highlight how architectural differences translate to real-world trading results:

Dimension	Multi-Output Strategy	Single-Output Strategy
Total Return	41.89% (test), 32.12% (OOS)	12.37% (test), -2.92% (OOS)
Risk-Adjusted Returns	Sharpe 0.45 (test), 1.86 (OOS)	Sharpe 0.16 (test), -0.12 (OOS)
Drawdown Resilience	-19.73% (test), ~-10%–-15% (OOS, inferred)	-27.63% (test), N/A (OOS not reported)
Consistency	Outperforms benchmark in both test/OOS periods	Underperforms benchmark in both periods

Why the Multi-Output Strategy Dominates

Trading strategies care more about ranking assets than predicting exact returns. Multi-output models are better at this because they learn from all assets together. Single-output models give isolated predictions without context (Qian, et al, 2022). For example, during the 2020 COVID-19 crash, multi-output models ranked GLD (gold) as the best and DBO (oil) as the worst, using cross-asset information. Single-output models failed to do this and made bad trades.

Market-neutral strategies with 100% long and 100% short positions can reduce overall market risk. Neutral middle assets avoid low-confidence trades. Single-output strategies still perform poorly because their rankings are less reliable.

The multi-output strategy's high total return (45.13% in testing) easily covered frequent rebalancing costs (-3.24%). The single-output strategy had moderate returns (15.49%), so similar costs reduced profits significantly.

When the Single-Output Strategy Might Narrow the Gap

In markets where assets move differently, single-output models can do better. For example, if gold goes up while stocks and bonds go down, the asset-specific focus helps. But this situation is rare. Most shocks, like rate hikes or recessions, affect many assets at the same time.

For small portfolios with high costs, single-output strategies are not very good. Rebalancing costs above 0.2% reduce returns, and their total returns are lower. This makes them less competitive compared to multi-output strategies.

Backtesting Performance Takeaway

The advantage of multi-output strategies in backtesting stems from their profit model aligning with how trading strategies operate: relative returns, rather than absolute predictability. The poor performance of single-output strategies reflects their failure to fully leverage cross-asset environments—although they outperform multi-output models slightly on individual assets such as GLD (Kara, et al, 2025).

3. Information Capture: Complementary but Non-Overlapping

A key question is whether the single-output and multi-output LSTM models capture the same information. The answer is no. Their design makes them capture different market dynamics.

Single-output LSTMs use only one asset's historical returns to predict its future returns. They learn asset-specific time patterns. For example, short-term momentum, seasonality, or mean reversion. But they do not see other assets. For example, a GLD model can learn that gold rises after three down days. It cannot see that gold may rise when SPY falls a lot.

Multi-output LSTMs use all five assets' returns as input. They capture shared market trends and cross-asset relationships. For example, if SPY, TLT, and DBO fall together, GLD may rise. They also learn conditional rules, like TLT up and SPY down predicts SHY will outperform DBO. But the shared LSTM weakens single-asset signals, so it does worse for assets like GLD compared to single-output models.

The overlap between the two models is small. Both capture some time patterns for specific assets, but single-output models do it better. For example, SPY momentum is learned more accurately by the single-output SPY model. Neither model sees external factors like inflation or geopolitical events. This limits both, but it affects multi-output models more because cross-asset relationships can break under shocks.

In summary, single-output and multi-output models capture different information. Single-output models focus on individual asset patterns. Multi-output models focus on cross-asset dependencies and market trends. This is why their trading results differ. Multi-output models are better for portfolio management because markets are connected.

The main insight is that architecture matters more than predicting a single asset perfectly. Single-output models predict one asset, but trading needs relative rankings. Multi-output models capture information across assets and give more consistent rankings. This makes backtests more successful, even if individual asset predictions have low R^2.

Key Implications for Practice

Choose a multi-output model if you want a portfolio-level strategy because it can use cross-asset information, even if it is weaker for single assets. Use a single-output model for a concentrated strategy with only one or two assets, like GLD, because it focuses on that asset and can perform better. A hybrid approach can also work well. You can use multi-output forecasts to rank assets and single-output forecasts to adjust position size, giving more weight to assets where the single-output model works best.

Final Thought

Neither architecture can "solve" the problem of financial predictability—markets are too complex to do so. However, multi-output models can capture cross-asset dependencies, making them more practical in real-world trading strategies. Their successful backtesting demonstrates that context is more important than expertise in a single asset in portfolio management. The inefficiency of single-output models reminds us that financial markets are not collections of isolated assets, but rather interconnected systems; models that ignore these connections struggle to generate consistent and stable returns.

FootNote

Advantage for using AI: Needing less time to collect information and to evaluate it for accuracy as compared to using bibliographical research.

References:

Archakov, Ilya, and Peter Reinhard Hansen. "A new parametrization of correlation matrices." *Econometrica* 89, no. 4 (2021): 1699-1715.

Huang, Ruijie, Chenji Wei, Baohua Wang, Jian Yang, Xin Xu, Suwei Wu, and Suqi Huang. "Well performance prediction based on Long Short-Term Memory (LSTM) neural network." *Journal of Petroleum Science and Engineering* 208 (2022): 109686.

Hortúa, Héctor J., Andrés Mora-Valencia, and Esteban Vanegas. "Stock Return Predictability on S&P 500 ETF: A Deep Learning Approach." *Annals of Data Science* (2025): 1-64.

Cipiloglu Yildiz, Zeynep, and Selim Baha Yildiz. "A portfolio construction framework using LSTM-based stock markets forecasting." *International Journal of Finance & Economics* 27, no. 2 (2022): 2356-2366.

Wei, Ze-Lin, Hong-Yu An, Yao Yao, Wei-Cong Su, Guo Li, Saifullah, Bi-Feng Sun, and Mu-Jiang-Shan Wang. "FSTGAT: Financial Spatio-Temporal Graph Attention Network for Non-Stationary Financial Systems and Its Application in Stock Price Prediction." *Symmetry* 17, no. 8 (2025): 1344.

Qian, Jie, and Gonggui Chen. "Improved Multi-goal Particle Swarm Optimization Algorithm and Multi-output BP Network for Optimal Operation of Power System." *IAENG International Journal of Applied Mathematics* 52, no. 3 (2022).

Kara, Hasan Bedri Taha, Alperen Sayar, Peri Güneş, Mert Güvençli, Seyit Ertuğrul, and Tuna Çakar. "Multi-Output vs Single-Output Deep Learning for Plant Disease Detection." In 2025 *10th International Conference on Computer Science and Engineering* (UBMK), pp. 190-195. IEEE, 2025.