ML Group Work Project 1 M3 Submission / GWP1_Machine_Learning_in_Finance - Colab.pdf

WorldQuant University

Document Details

Submission ID

trn:oid:::8628:502389983

Submission Date

Sep 23, 2025, 3:51 PM GMT+3

Download Date

Sep 23, 2025, 3:54 PM GMT+3

File Name

GWP1_Machine_Learning_in_Finance - Colab.pdf

File Size

1018.4 KB

19 Pages

3,700 Words

20,980 Characters

34% Overall Similarity

The combined total of all matches, including overlapping sources, for each database.

Filtered from the Report

- Bibliography
- Small Matches (less than 30 words)

Exclusions

- 1 Excluded Source

Match Groups

16 Not Cited or Quoted 34%
Matches with neither in-text citation nor quotation marks

0 Missing Quotations 0%
Matches that are still very similar to source material

0 Missing Citation 0%
Matches that have quotation marks, but no in-text citation

0 Cited and Quoted 0%
Matches with in-text citation present, but no quotation marks

Top Sources

21% Internet sources
4% Publications
32% Submitted works (Student Papers)

Integrity Flags

0 Integrity Flags for Review

Our system's algorithms look deeply at a document for any inconsistencies that would set it apart from a normal submission. If we notice something strange, we flag it for you to review.

A Flag is not necessarily an indicator of a problem. However, we'd recommend you focus your attention there for further review.

Match Groups

- **16** Not Cited or Quoted 34%
  Matches with neither in-text citation nor quotation marks
- **0** Missing Quotations 0%
  Matches that are still very similar to source material
- **0** Missing Citation 0%
  Matches that have quotation marks, but no in-text citation
- **0** Cited and Quoted 0%
  Matches with in-text citation present, but no quotation marks

Top Sources

21% Internet sources
4% Publications
32% Submitted works (Student Papers)

Top Sources

The sources with the highest number of matches within the submission. Overlapping sources will not be displayed.

1	Submitted works	University of North Texas on 2025-09-18	13%
2	Submitted works	University of North Texas on 2022-10-30	5%
3	Submitted works	Glasgow Caledonian University on 2024-01-03	3%
4	Internet	medium.com	2%
5	Submitted works	University of Greenwich on 2025-04-03	2%
6	Internet	www.digitalocean.com	1%
7	Submitted works	Glasgow Caledonian University on 2023-12-18	1%
8	Submitted works	University of North Texas on 2023-11-03	1%
9	Submitted works	Colorado Technical University Online on 2025-07-12	1%
10	Publication	Jan Y. Yang. "Chapter 3 Case Studies: AI in Action", Springer Science and Business ...	1%

11 Submitted works

Griffth University on 2024-10-07 <1%

12 Internet

www.codemotion.com <1%

13 Submitted works

The Robert Gordon University on 2025-07-07 <1%

14 Submitted works

Swinburne University of Technology on 2024-10-18 <1%

15 Submitted works

Glasgow Caledonian University on 2024-01-15 <1%

Category 1:Elastic nets

```
import numpy as np
import pandas as pd
from sklearn.linear_model import ElasticNetCV
from sklearn.datasets import make_regression
from sklearn.model_selection import train_test_split

# Simulated data
X, y = make_regression(n_samples=500, n_features=20, noise=0.1, random_state=42)
# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Elastic Net with cross-validation
elastic = ElasticNetCV(l1_ratio=[.1, .5, .7, .9, .95, .99, 1],  # grid of alpha
                       alphas=np.logspace(-3, 1, 50),            # grid of lambdas
                       cv=5,
                       random_state=42)

elastic.fit(X_train, y_train)
print("Best alpha :", elastic.alpha_)
print("Best l1_ratio (α):", elastic.l1_ratio_)
print("R^2 on test:", elastic.score(X_test, y_test))
```

```
Best alpha : 0.0021209508879201904
Best l1_ratio (α): 1.0
R^2 on test: 0.9999994275069372
```

CATEGORY 2: K-Means Clustering

```
!pip install yfinance
```

```
Requirement already satisfied: yfinance in /usr/local/lib/python3.12/dist-packages (0.2.65)
Requirement already satisfied: pandas>=1.3.0 in /usr/local/lib/python3.12/dist-packages (from yfinance) (2.2.2)
Requirement already satisfied: numpy>=1.16.5 in /usr/local/lib/python3.12/dist-packages (from yfinance) (2.0.2)
Requirement already satisfied: requests>=2.31 in /usr/local/lib/python3.12/dist-packages (from yfinance) (2.32.4)
Requirement already satisfied: multitasking>=0.0.7 in /usr/local/lib/python3.12/dist-packages (from yfinance) (0.0.12)
Requirement already satisfied: platformdirs>=2.0.0 in /usr/local/lib/python3.12/dist-packages (from yfinance) (4.4.0)
Requirement already satisfied: pytz>=2022.5 in /usr/local/lib/python3.12/dist-packages (from yfinance) (2025.2)
Requirement already satisfied: frozendict>=2.3.4 in /usr/local/lib/python3.12/dist-packages (from yfinance) (2.4.6)
Requirement already satisfied: peewee>=3.16.2 in /usr/local/lib/python3.12/dist-packages (from yfinance) (3.18.2)
Requirement already satisfied: beautifulsoup4>=4.11.1 in /usr/local/lib/python3.12/dist-packages (from yfinance) (4.13.5)
Requirement already satisfied: curl_cffi>=0.7 in /usr/local/lib/python3.12/dist-packages (from yfinance) (0.13.0)
Requirement already satisfied: protobuf>=3.19.0 in /usr/local/lib/python3.12/dist-packages (from yfinance) (5.29.5)
Requirement already satisfied: websockets>=13.0 in /usr/local/lib/python3.12/dist-packages (from yfinance) (15.0.1)
Requirement already satisfied: soupsieve>1.2 in /usr/local/lib/python3.12/dist-packages (from beautifulsoup4>=4.11.1->yfinance) (2.8)
```

turnitin Page 6 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

```
Requirement already satisfied: typing-extensions>=4.0.0 in /usr/local/lib/python3.12/dist-packages (from beautifulsoup4>=4.11.1->yfinance) (4.15.0)
Requirement already satisfied: cffi>=1.12.0 in /usr/local/lib/python3.12/dist-packages (from curl_cffi>=0.7->yfinance) (2.0.0)
Requirement already satisfied: certifi>=2024.2.2 in /usr/local/lib/python3.12/dist-packages (from curl_cffi>=0.7->yfinance) (2025.8.3)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas>=1.3.0->yfinance) (2.9.0.post0)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas>=1.3.0->yfinance) (2025.2)
Requirement already satisfied: charset_normalizer<4,>=2 in /usr/local/lib/python3.12/dist-packages (from requests>=2.31->yfinance) (3.4.3)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.12/dist-packages (from requests>=2.31->yfinance) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.12/dist-packages (from requests>=2.31->yfinance) (2.5.0)
Requirement already satisfied: pycparser in /usr/local/lib/python3.12/dist-packages (from cffi>=1.12.0->curl_cffi>=0.7->yfinance) (2.23)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.8.2->pandas>=1.3.0->yfinance) (1.17.0)
```

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
from sklearn.decomposition import PCA
from scipy.cluster.hierarchy import linkage, dendrogram, fcluster
from scipy.spatial.distance import squareform
import yfinance as yf
import warnings
```

```
warnings.filterwarnings("ignore")
sns.set(style="whitegrid")
```

```
tickers = [
    "AAPL","MSFT","GOOGL","AMZN","META","JPM","BAC","WMT","PG","KO",
    "XOM","CVX","TSLA","NVDA","INTC","CSCO","ADBE","ORCL","NKE","MCD"
]  # sample 20 US large caps
lookback_days = 252 * 2  # two years of daily returns
end_date = None  # today
start_date = None
```

```
# download prices and compute log returns
data = yf.download(tickers, period="3y", interval="1d", progress=False, auto_adjust=False )["Adj Close"]
data = data.dropna(axis=1, how='any')  # drop tickers with missing data
prices = data.ffill().dropna()
returns = np.log(prices).diff().dropna()
# Trim to lookback
returns = returns.tail(lookback_days)
returns.shape
```

```
(504, 20)
```

```
# correlation and correlation distance
corr = returns.corr()  # Pearson correlation
# correlation distance as commonly used:
```

turnitin Page 6 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

turnitin Page 7 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

```
dist_corr = np.sqrt(0.5 * (1 - corr))
# convert to condensed distance matrix for scipy hierarchical methods
condensed_dist = squareform(dist_corr.values, checks=False)
```

```
# hierarchical clustering (Ward) and dendrogram
Z = linkage(condensed_dist, method="ward")  # ward on distance matrix
plt.figure(figsize=(6, 4))
dendro = dendrogram(Z, labels=corr.index.tolist(), leaf_rotation=90)
plt.title("Hierarchical Clustering Dendrogram (Ward, correlation-distance)")
plt.tight_layout()
plt.show()
```

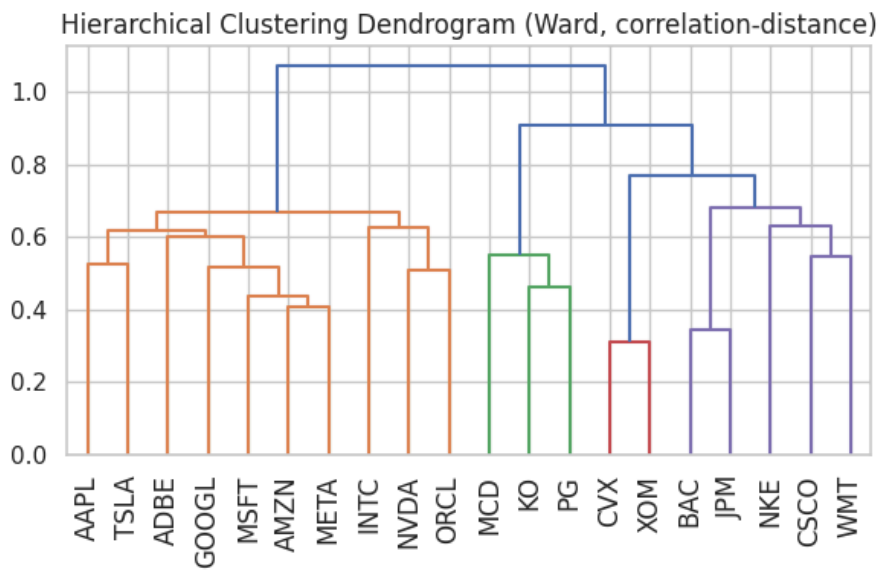


```
# choose number of clusters from dendrogram/ specify K
K_from_dendrogram = 5
labels_hier = fcluster(Z, K_from_dendrogram, criterion="maxclust")
cluster_map_hier = pd.Series(labels_hier, index=corr.index, name="cluster_hier")
print(cluster_map_hier.value_counts().sort_index())
```

```
cluster_hier
1    10
2     3
3     2
4     2
5     3
Name: count, dtype: int64
```

turnitin Page 7 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

turnitin Page 8 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

```
# clustered correlation heatmap
order = dendro['leaves']  # order of tickers from dendrogram
ordered_tickers = [corr.index[i] for i in order]
corr_ordered = corr.loc[ordered_tickers, ordered_tickers]

plt.figure(figsize=(5, 3))
sns.heatmap(corr_ordered, cmap="vlag", center=0, xticklabels=True, yticklabels=True)
plt.title("Clustered Correlation Heatmap (ordered by dendrogram)")
plt.show()
```

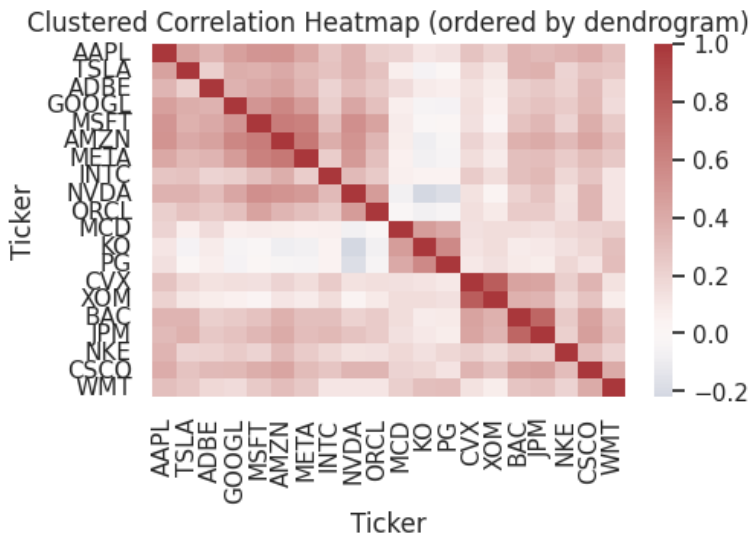


```
# k-means on PCA-reduced returns (common pipeline)
# Standard approach: reduce dimensionality (PCA) then kmeans on leading PCs
pca = PCA(n_components=5)
X_pca = pca.fit_transform(returns.T)  # assets x features -> shape (n_assets, n_components)

# iterate K candidates using silhouette
sil_scores = {}
# The number of clusters must be less than the number of samples (tickers).
K_range = range(2, len(returns.T) -1 )
for K in K_range:
    km = KMeans(n_clusters=K, n_init=50, random_state=42)
    labels = km.fit_predict(X_pca)
    sil = silhouette_score(X_pca, labels)
    sil_scores[K] = sil

# plot silhouette vs K
plt.figure(figsize=(6,3))
plt.plot(list(sil_scores.keys()), list(sil_scores.values()), marker='o')
plt.xlabel("K")
```

turnitin Page 8 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

turnitin Page 9 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

```
plt.ylabel("Silhouette score")
plt.title("Silhouette score vs K (k-means on PCA features)")
plt.show()

best_K = max(sil_scores, key=sil_scores.get)
print("best K by silhouette:", best_K)
```

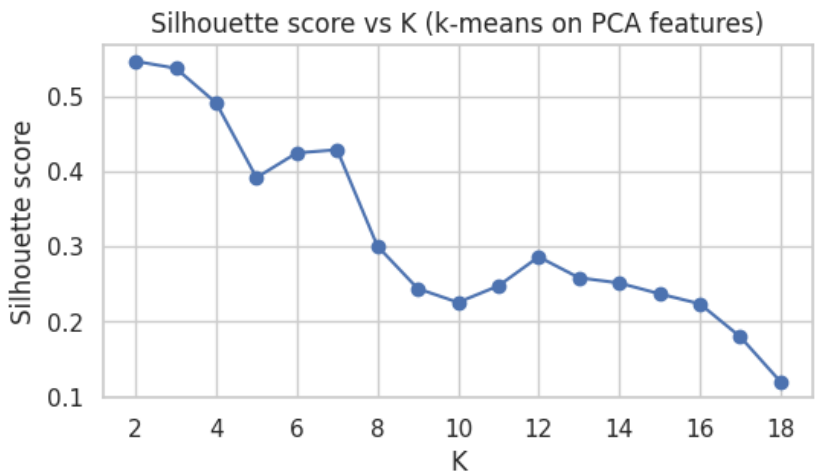


```
best K by silhouette: 2
```

```
# k-means clustering and cluster summary
km = KMeans(n_clusters=best_K, n_init=100, random_state=42)
labels_km = km.fit_predict(X_pca)
cluster_map_km = pd.Series(labels_km+1, index=returns.columns, name="cluster_kmeans")  # +1 to make 1..K
cluster_map_km.value_counts().sort_index()
```

cluster_kmeans	count
1	19
2	1

dtype: int64

```
# 2D projection of assets by first 2 PCs colored by cluster
plt.figure(figsize=(5,3))
for c in np.unique(labels_km):
    idx = np.where(labels_km==c)
    plt.scatter(X_pca[idx,0], X_pca[idx,1], label=f"Cluster {c+1}", s=80, alpha=0.8)
plt.xlabel("PC1"); plt.ylabel("PC2")
plt.title("Assets in PCA space colored by k-means cluster")
plt.legend()
```

turnitin Page 9 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

turnitin Page 10 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

```
plt.legend()
plt.show()
```

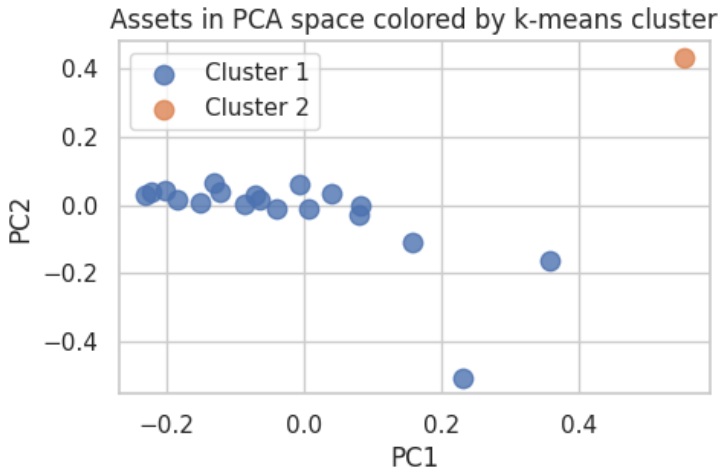


```
# cluster-level return/volatility statistics (on historical returns)
summary = []
for cl in sorted(cluster_map_km.unique()):
    assets = cluster_map_km[cluster_map_km==cl].index
    r = returns[assets].mean(axis=1)
    summary.append({
        "cluster": cl,
        "n_assets": len(assets),
        "avg_daily_ret": r.mean(),
        "ann_ret": r.mean()*252,
        "ann_vol": r.std()*np.sqrt(252),
        "sharpe": (r.mean()*252)/(r.std()*np.sqrt(252)+1e-9)
    })
summary_df = pd.DataFrame(summary).set_index("cluster")
summary_df
```

cluster	n_assets	avg_daily_ret	ann_ret	ann_vol	sharpe
1	19	0.000756	0.190441	0.161396	1.179962
2	1	0.000873	0.220027	0.624720	0.352201

CATEGORY 4: Classification Trees

turnitin Page 10 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

```
# Classification Trees - Jupyter Notebook Example
```

turnitin Page 11 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

```
# ----------------------------------------------
# This notebook illustrates how a classification tree works
# using the Iris dataset as an example.

# Step 1: Import libraries
import numpy as np
import pandas as pd
from sklearn.datasets import load_iris
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, accuracy_score
import matplotlib.pyplot as plt

# Step 2: Load the dataset
iris = load_iris()
X = iris.data
y = iris.target
feature_names = iris.feature_names
target_names = iris.target_names

# Convert to DataFrame for better readability
df = pd.DataFrame(X, columns=feature_names)
df['target'] = y
df.head()

# Step 3: Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42
)

# Step 4: Train the classification tree
clf = DecisionTreeClassifier(criterion="gini", max_depth=3, random_state=42)
clf.fit(X_train, y_train)

# Step 5: Visualize the classification tree
plt.figure(figsize=(12,8))
plot_tree(clf,
          feature_names=feature_names,
          class_names=target_names,
          filled=True,
          rounded=True)
plt.show()

# Step 6: Make predictions
y_pred = clf.predict(X_test)

# Step 7: Evaluate performance
print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n")
print(classification_report(y_test, y_pred, target_names=target_names))
```

turnitin Page 11 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

turnitin Page 12 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

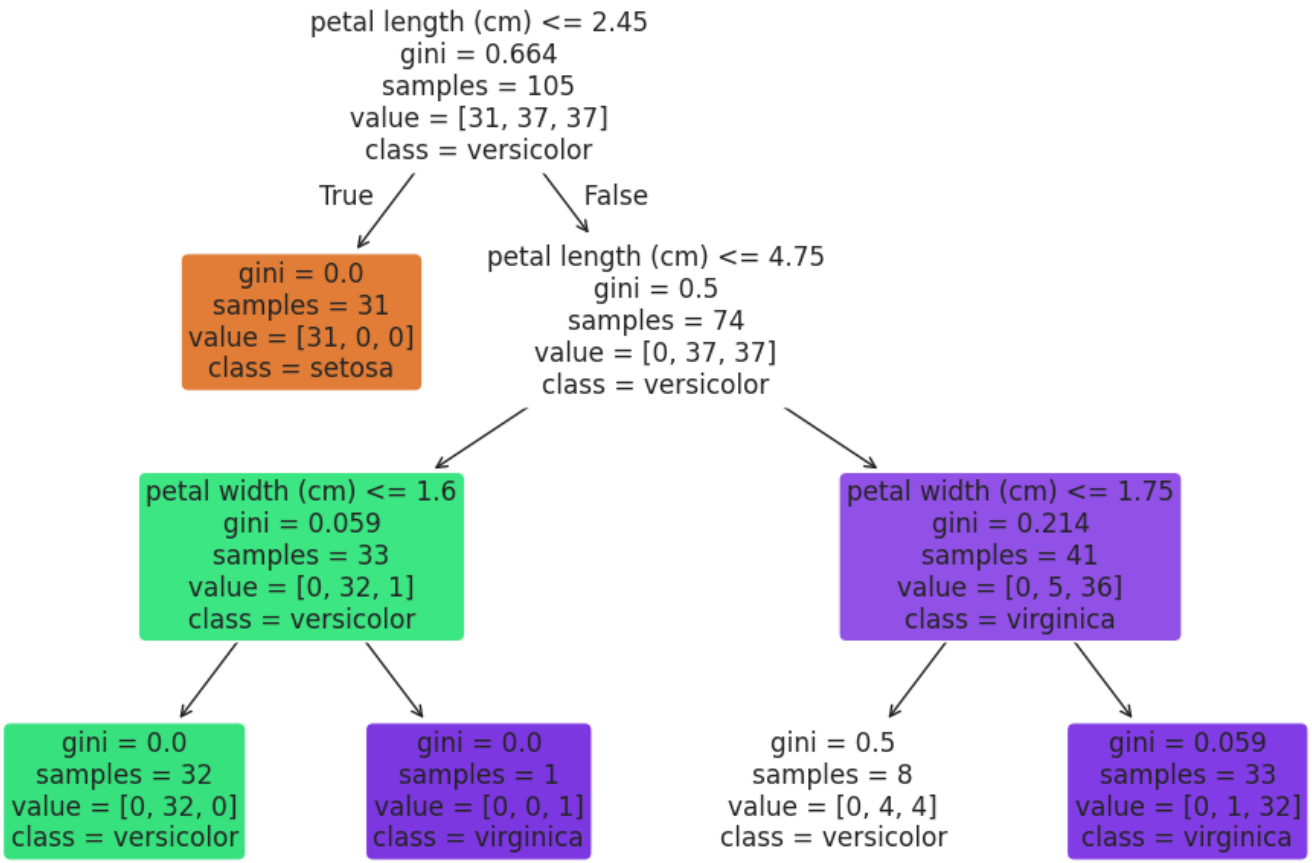


```
Accuracy: 1.0

Classification Report:

              precision    recall  f1-score   support

      setosa       1.00      1.00      1.00        19
  versicolor       1.00      1.00      1.00        13
   virginica       1.00      1.00      1.00        13

    accuracy                           1.00        45
   macro avg       1.00      1.00      1.00        45
weighted avg       1.00      1.00      1.00        45
```

```
# Customer Churn Prediction using Classification Trees
# ----------------------------------------------------
# This notebook demonstrates how classification trees can be applied
```

turnitin Page 12 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

turnitin Page 13 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

```
# in a marketing context (predicting whether a customer will churn).

# Step 1: Import libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import classification_report, accuracy_score

# Step 2: Load dataset
# If you have the Kaggle dataset, replace the path below:
# df = pd.read_csv("Telco-Customer-Churn.csv")

# For demonstration, we'll create a smaller synthetic dataset
data = {
    "Customer_Age": [25, 45, 33, 50, 29, 62, 41, 37, 23, 48],
    "Monthly_Spend": [40, 90, 65, 120, 50, 130, 80, 70, 35, 110],
    "Tenure_Months": [3, 24, 12, 36, 6, 48, 20, 18, 2, 30],
    "Support_Calls": [1, 4, 2, 5, 1, 6, 3, 2, 0, 4],
    "Churn": [1, 0, 1, 0, 1, 0, 0, 1, 1, 0]  # 1 = churn, 0 = stayed
}
df = pd.DataFrame(data)

print("Sample Data:")
print(df.head())

# Step 3: Features and target
X = df.drop("Churn", axis=1)
y = df["Churn"]

# Step 4: Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42
)

# Step 5: Train classification tree
clf = DecisionTreeClassifier(criterion="entropy", max_depth=3, random_state=42)
clf.fit(X_train, y_train)

# Step 6: Visualize the tree
plt.figure(figsize=(10,6))
plot_tree(clf,
          feature_names=X.columns,
          class_names=["Stayed", "Churned"],
          filled=True,
          rounded=True)
plt.show()

# Step 7: Predictions
y_pred = clf.predict(X_test)
```

turnitin Page 13 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

```
# Step 8: Evaluation
print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n")
print(classification_report(y_test, y_pred, target_names=["Stayed", "Churned"]))
```

```
Sample Data:
   Customer_Age  Monthly_Spend  Tenure_Months  Support_Calls  Churn
0            25             40              3              1      1
1            45             90             24              4      0
2            33             65             12              2      1
3            50            120             36              5      0
4            29             50              6              1      1
```

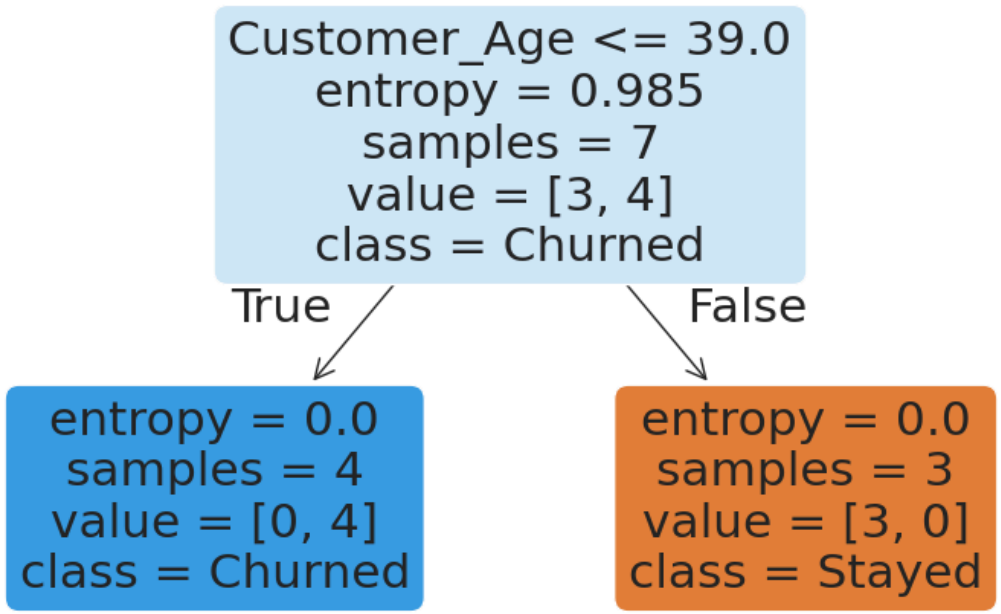


```
Accuracy: 1.0

Classification Report:

              precision    recall  f1-score   support

      Stayed       1.00      1.00      1.00         2
     Churned       1.00      1.00      1.00         1

    accuracy                           1.00         3
   macro avg       1.00      1.00      1.00         3
weighted avg       1.00      1.00      1.00         3
```

turnitin Page 15 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

```
# Customer Churn Prediction using Classification Trees
# ----------------------------------------------------
# This notebook demonstrates how classification trees can be applied
# in a marketing context (predicting whether a customer will churn).

# Step 1: Import libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import classification_report, accuracy_score
import os

# Step 2: Try downloading dataset from Kaggle using kagglehub
try:
    import kagglehub
    path = kagglehub.dataset_download("blastchar/telco-customer-churn")
    dataset_file = os.path.join(path, "WA_Fn-UseC_-Telco-Customer-Churn.csv")
    print("Downloaded dataset via KaggleHub.")
except Exception as e:
    print("KaggleHub download failed:", e)
    # Fallback: look for local CSV file in working directory
    dataset_file = "WA_Fn-UseC_-Telco-Customer-Churn.csv"
    if os.path.exists(dataset_file):
        print("Loaded dataset from local CSV file.")
    else:
        raise FileNotFoundError(
            "Dataset not found. Please place 'WA_Fn-UseC_-Telco-Customer-Churn.csv' "
            "in your working directory."
        )

# Step 3: Load the dataset
df = pd.read_csv(dataset_file)

print("Dataset shape:", df.shape)
print(df.head())

# Step 4: Data preprocessing
# Drop customerID since it's just an identifier
df = df.drop("customerID", axis=1)

# Convert TotalCharges to numeric (it may contain blanks)
df["TotalCharges"] = pd.to_numeric(df["TotalCharges"], errors="coerce")
df["TotalCharges"] = df["TotalCharges"].fillna(df["TotalCharges"].median())

# Encode target variable: Churn (Yes=1, No=0)
df["Churn"] = df["Churn"].map({"Yes": 1, "No": 0})

# One-hot encode categorical variables
```

turnitin Page 15 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

turnitin Page 16 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

```
df_encoded = pd.get_dummies(df, drop_first=True)

# Step 5: Split into features (X) and target (y)
X = df_encoded.drop("Churn", axis=1)
y = df_encoded["Churn"]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42, stratify=y
)

# Step 6: Train a classification tree
clf = DecisionTreeClassifier(criterion="entropy", max_depth=4, random_state=42)
clf.fit(X_train, y_train)

# Step 7: Visualize the decision tree
plt.figure(figsize=(20,10))
plot_tree(clf,
          feature_names=X.columns,
          class_names=["Stayed", "Churned"],
          filled=True,
          rounded=True,
          fontsize=8)
plt.show()

# Step 8: Evaluate model performance
y_pred = clf.predict(X_test)

print("Accuracy on test set:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n")
print(classification_report(y_test, y_pred, target_names=["Stayed", "Churned"]))

# Step 9: Feature importance (marketing insight)
feature_importance = pd.DataFrame({
    "Feature": X.columns,
    "Importance": clf.feature_importances_
}).sort_values(by="Importance", ascending=False)

print("\nTop 10 Features Influencing Churn:\n")
print(feature_importance.head(10))

sns.barplot(x="Importance", y="Feature", data=feature_importance.head(10))
plt.title("Top Factors Influencing Customer Churn")
plt.show()
```

turnitin Page 16 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

turnitin Page 18 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

```
Using Colab cache for faster access to the 'telco-customer-churn' dataset.
Downloaded dataset via KaggleHub.
Dataset shape: (7043, 21)
   customerID  gender  SeniorCitizen Partner Dependents  tenure PhoneService  \
0  7590-VHVEG  Female              0     Yes         No       1           No
1  5575-GNVDE    Male              0      No         No      34          Yes
2  3668-QPYBK    Male              0      No         No       2          Yes
3  7795-CFOCW    Male              0      No         No      45           No
4  9237-HQITU  Female              0      No         No       2          Yes

      MultipleLines InternetService OnlineSecurity  ... DeviceProtection  \
0  No phone service             DSL             No  ...               No
1                No             DSL            Yes  ...              Yes
2                No             DSL            Yes  ...               No
3  No phone service             DSL            Yes  ...              Yes
4                No     Fiber optic             No  ...               No

  TechSupport StreamingTV StreamingMovies        Contract PaperlessBilling  \
0          No          No              No  Month-to-month              Yes
1          No          No              No        One year               No
2          No          No              No  Month-to-month              Yes
3         Yes          No              No        One year               No
4          No          No              No  Month-to-month              Yes

               PaymentMethod MonthlyCharges  TotalCharges Churn
0           Electronic check          29.85         29.85    No
1               Mailed check          56.95        1889.5    No
2               Mailed check          53.85        108.15   Yes
3  Bank transfer (automatic)          42.30       1840.75    No
4           Electronic check          70.70        151.65   Yes

[5 rows x 21 columns]
```

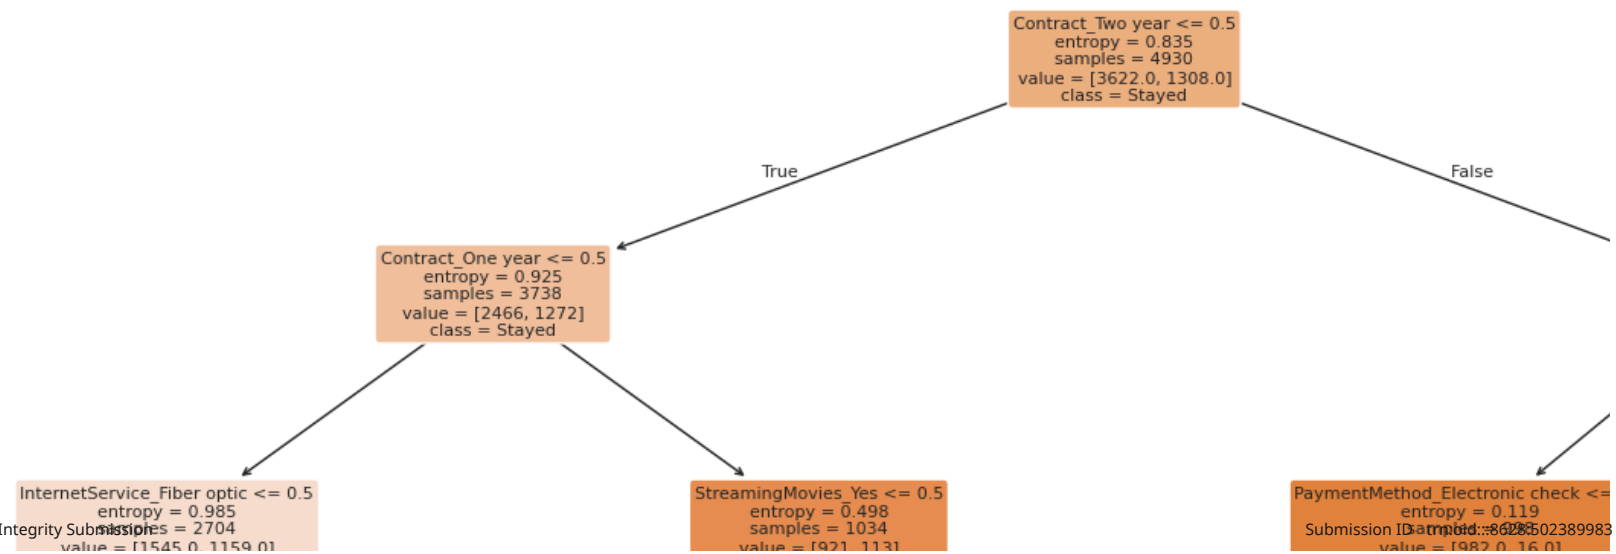

turnitin Page 18 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

turnitin Page 19 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

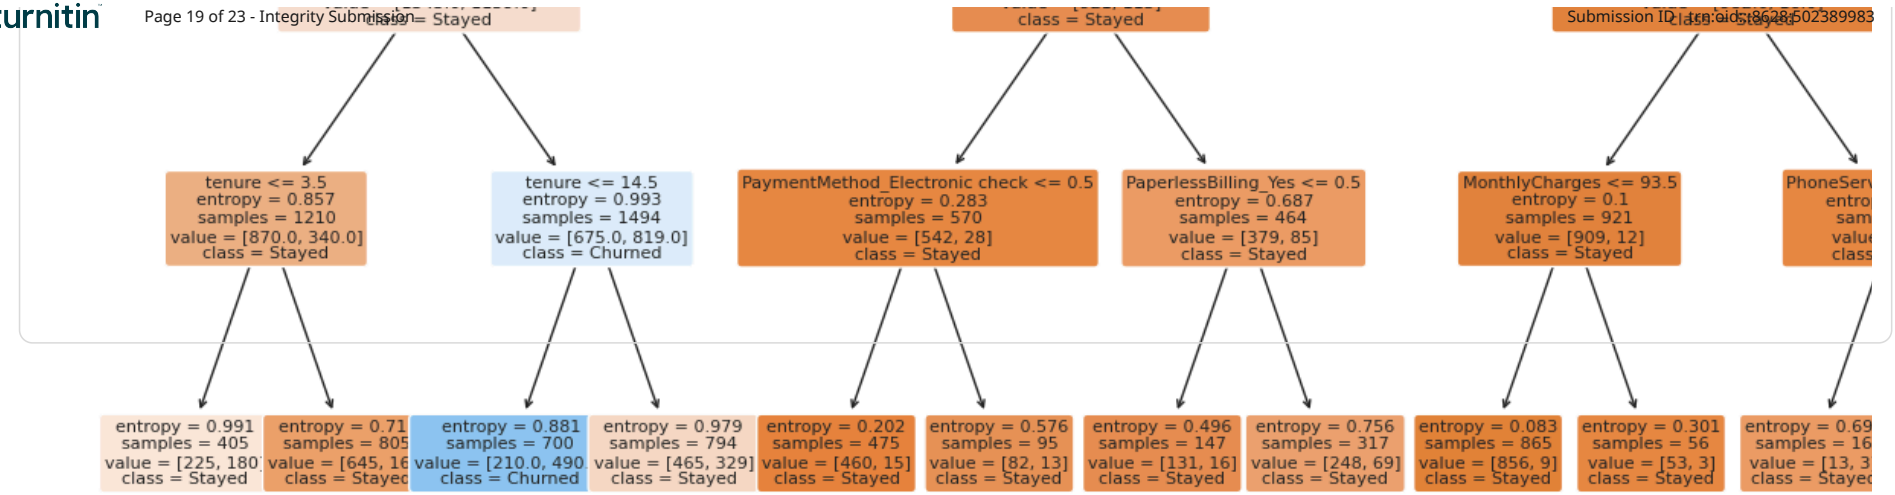


```
Accuracy on test set: 0.7875059157595835

Classification Report:

              precision    recall  f1-score   support

      Stayed       0.80      0.94      0.87      1552
     Churned       0.69      0.36      0.47       561

    accuracy                           0.79      2113
   macro avg       0.75      0.65      0.67      2113
weighted avg       0.77      0.79      0.76      2113


Top 10 Features Influencing Churn:

                         Feature  Importance
25              Contract_Two year    0.386092
24              Contract_One year    0.254277
1                          tenure    0.133098
10     InternetService_Fiber optic    0.130609
23            StreamingMovies_Yes    0.031684
2                  MonthlyCharges    0.022779
3                    TotalCharges    0.016857
28  PaymentMethod_Electronic check    0.012394
26            PaperlessBilling_Yes    0.005630
7                PhoneService_Yes    0.003811
```

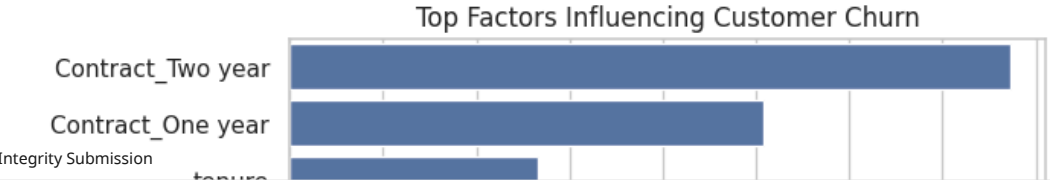

turnitin Page 19 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

turnitin Page 20 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

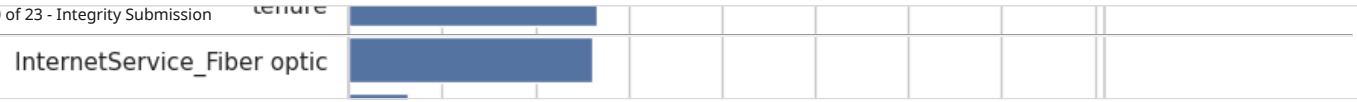


```
# Step 6: Hyperparameter tuning with GridSearchCV
from sklearn.model_selection import GridSearchCV

# Define parameter grid
param_grid = {
    "criterion": ["gini", "entropy", "log_loss"],
    "max_depth": [3, 5, 7, None],
    "min_samples_split": [2, 5, 10],
    "min_samples_leaf": [1, 2, 5],
    "max_features": [None, "sqrt", "log2"],
    "class_weight": [None, "balanced"]
}

# Initialize base model
dtree = DecisionTreeClassifier(random_state=42)

# Grid search with 5-fold cross-validation
grid_search = GridSearchCV(
    estimator=dtree,
    param_grid=param_grid,
    cv=5,
    scoring="accuracy",
    n_jobs=-1,
    verbose=1
)

grid_search.fit(X_train, y_train)

# Best hyperparameters
print("Best Parameters:", grid_search.best_params_)
print("Best CV Accuracy:", grid_search.best_score_)

# Step 7: Retrain with best parameters
best_dtree = grid_search.best_estimator_
best_dtree.fit(X_train, y_train)

# Step 8: Evaluate tuned model
y_pred_best = best_dtree.predict(X_test)

print("\nTest Set Accuracy:", accuracy_score(y_test, y_pred_best))
print("\nClassification Report (Tuned Model):\n")
print(classification_report(y_test, y_pred_best, target_names=["Stayed", "Churned"]))

# Step 9: Visualize tuned decision tree
plt.figure(figsize=(20,10))
plot_tree(best_dtree,
          feature_names=X_columns,
```

turnitin Page 20 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

turnitin Page 21 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

```
            class_names=["Stayed", "Churned"],
            filled=True,
            rounded=True,
            fontsize=8)
plt.show()

# Step 10: Feature importance from tuned model
feature_importance = pd.DataFrame({
    "Feature": X.columns,
    "Importance": best_dtree.feature_importances_
}).sort_values(by="Importance", ascending=False)

print("\nTop 10 Features Influencing Churn (Tuned Model):\n")
print(feature_importance.head(10))

sns.barplot(x="Importance", y="Feature", data=feature_importance.head(10))
plt.title("Top Factors Influencing Customer Churn (Tuned Tree)")
plt.show()
```

turnitin Page 21 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

turnitin Page 23 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983

```
Fitting 5 folds for each of 648 candidates, totalling 3240 fits
Best Parameters: {'class_weight': None, 'criterion': 'gini', 'max_depth': 5, 'max_features': None, 'min_samples_leaf': 5, 'min_samples_split': 2}
Best CV Accuracy: 0.7892494929006085

Test Set Accuracy: 0.7879791765262659

Classification Report (Tuned Model):

              precision    recall  f1-score   support

      Stayed       0.84      0.89      0.86      1552
     Churned       0.62      0.52      0.56       561

    accuracy                           0.79      2113
   macro avg       0.73      0.70      0.71      2113
weighted avg       0.78      0.79      0.78      2113
```

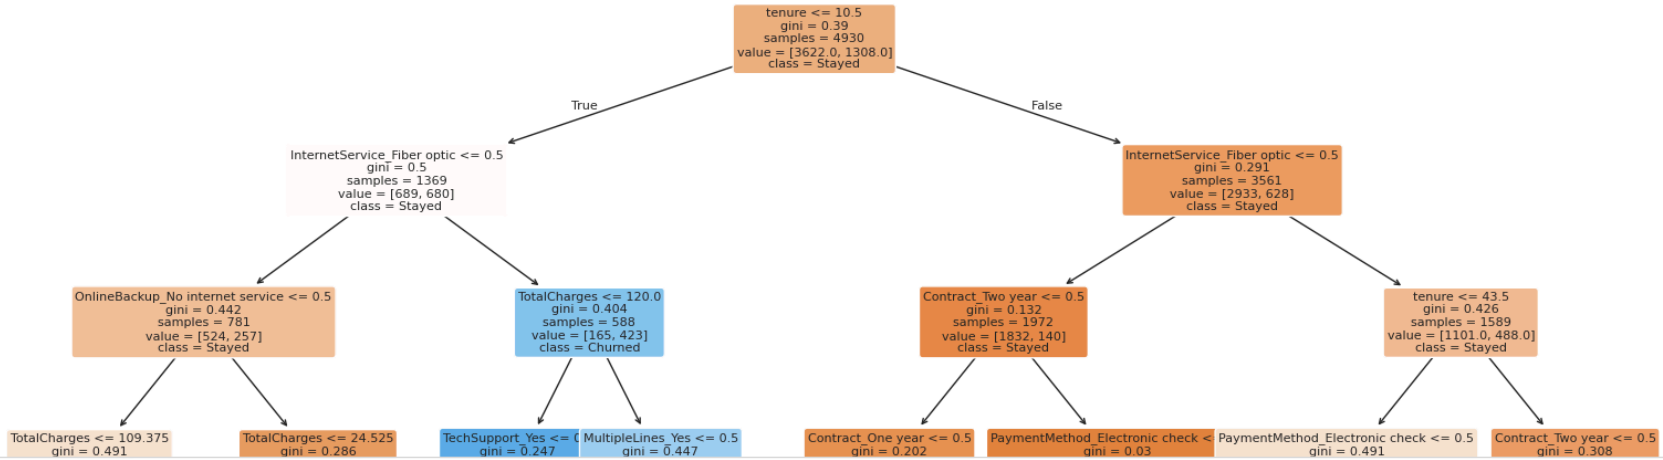


```
# ================================
# Step 6A: Hyperparameter tuning with GridSearchCV
# ================================
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from scipy.stats import randint

# Define parameter grid for GridSearchCV
param_grid = {
    "criterion": ["gini", "entropy", "log_loss"],
    "max_depth": [3, 5, 7, None],
    "min_samples_split": [2, 5, 10],
    "min_samples_leaf": [1, 2, 5],
    "max_features": [None, "sqrt", "log2"],
    "class_weight": [None, "balanced"]
}

dtree = DecisionTreeClassifier(random_state=42)
```

turnitin Page 23 of 23 - Integrity Submission Submission ID trn:oid:::8628:502389983